

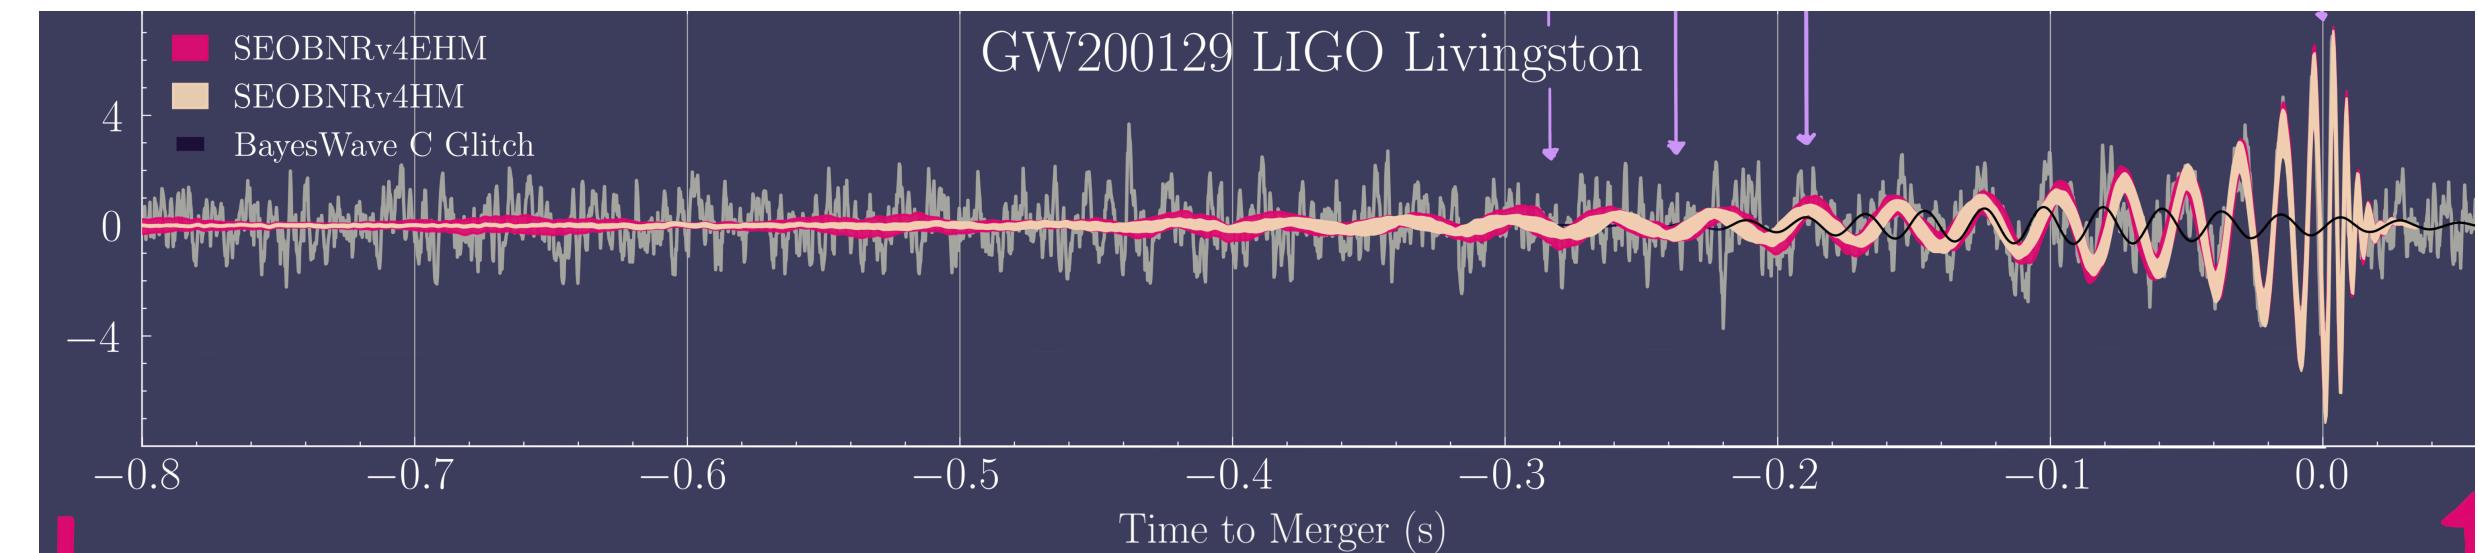
Overview on Deep Generative Modeling in the Natural Sciences

Introduction

Annalena Kofler, 16.04.2024

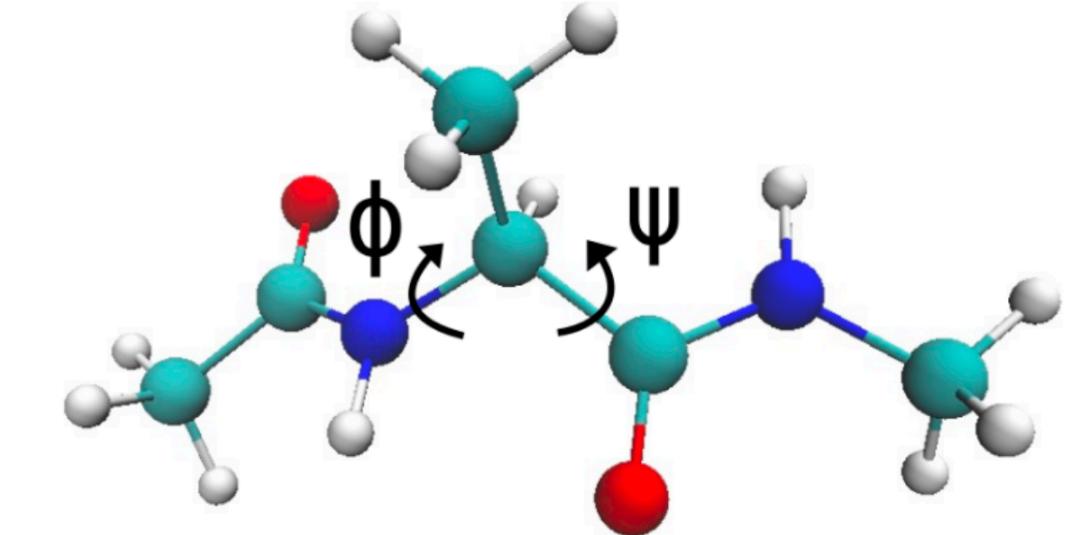
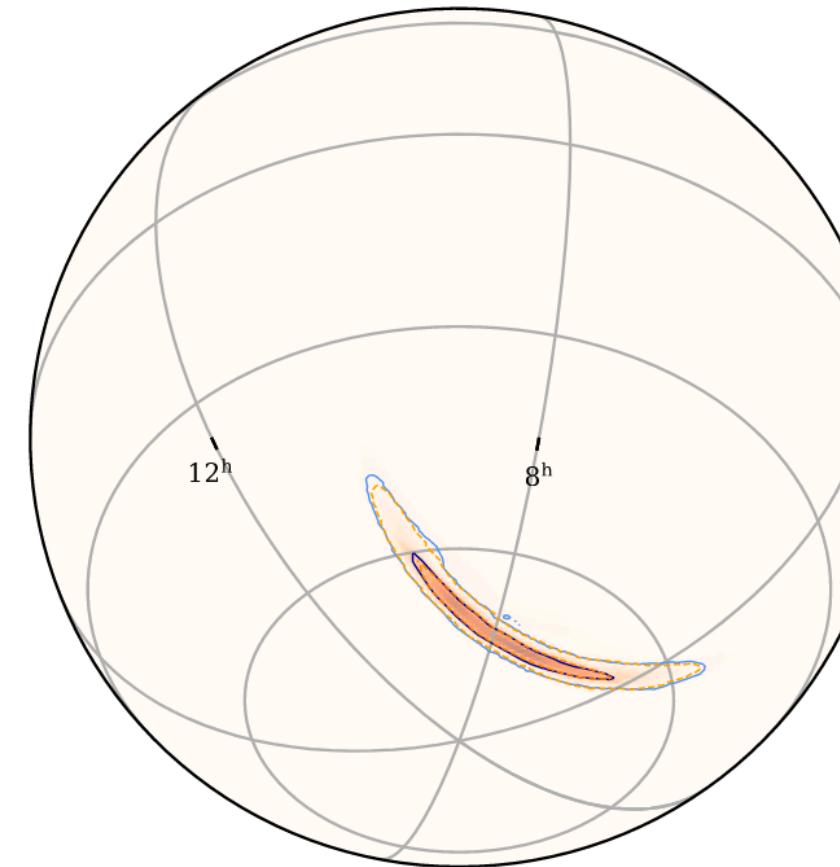
Short Introduction: Who am I?

- PhD student at the Max-Planck-Institute for Intelligent Systems in Tübingen supervised by Prof. Bernhard Schölkopf
- Working on ML for Physics
 - Posterior Estimation for Gravitational Wave Inference
 - Deep Generative Modeling & Simulation-based inference
 - Differentiable Physics Simulators
- Previously: Master's student at TUM working with Prof. Lukas Heinrich

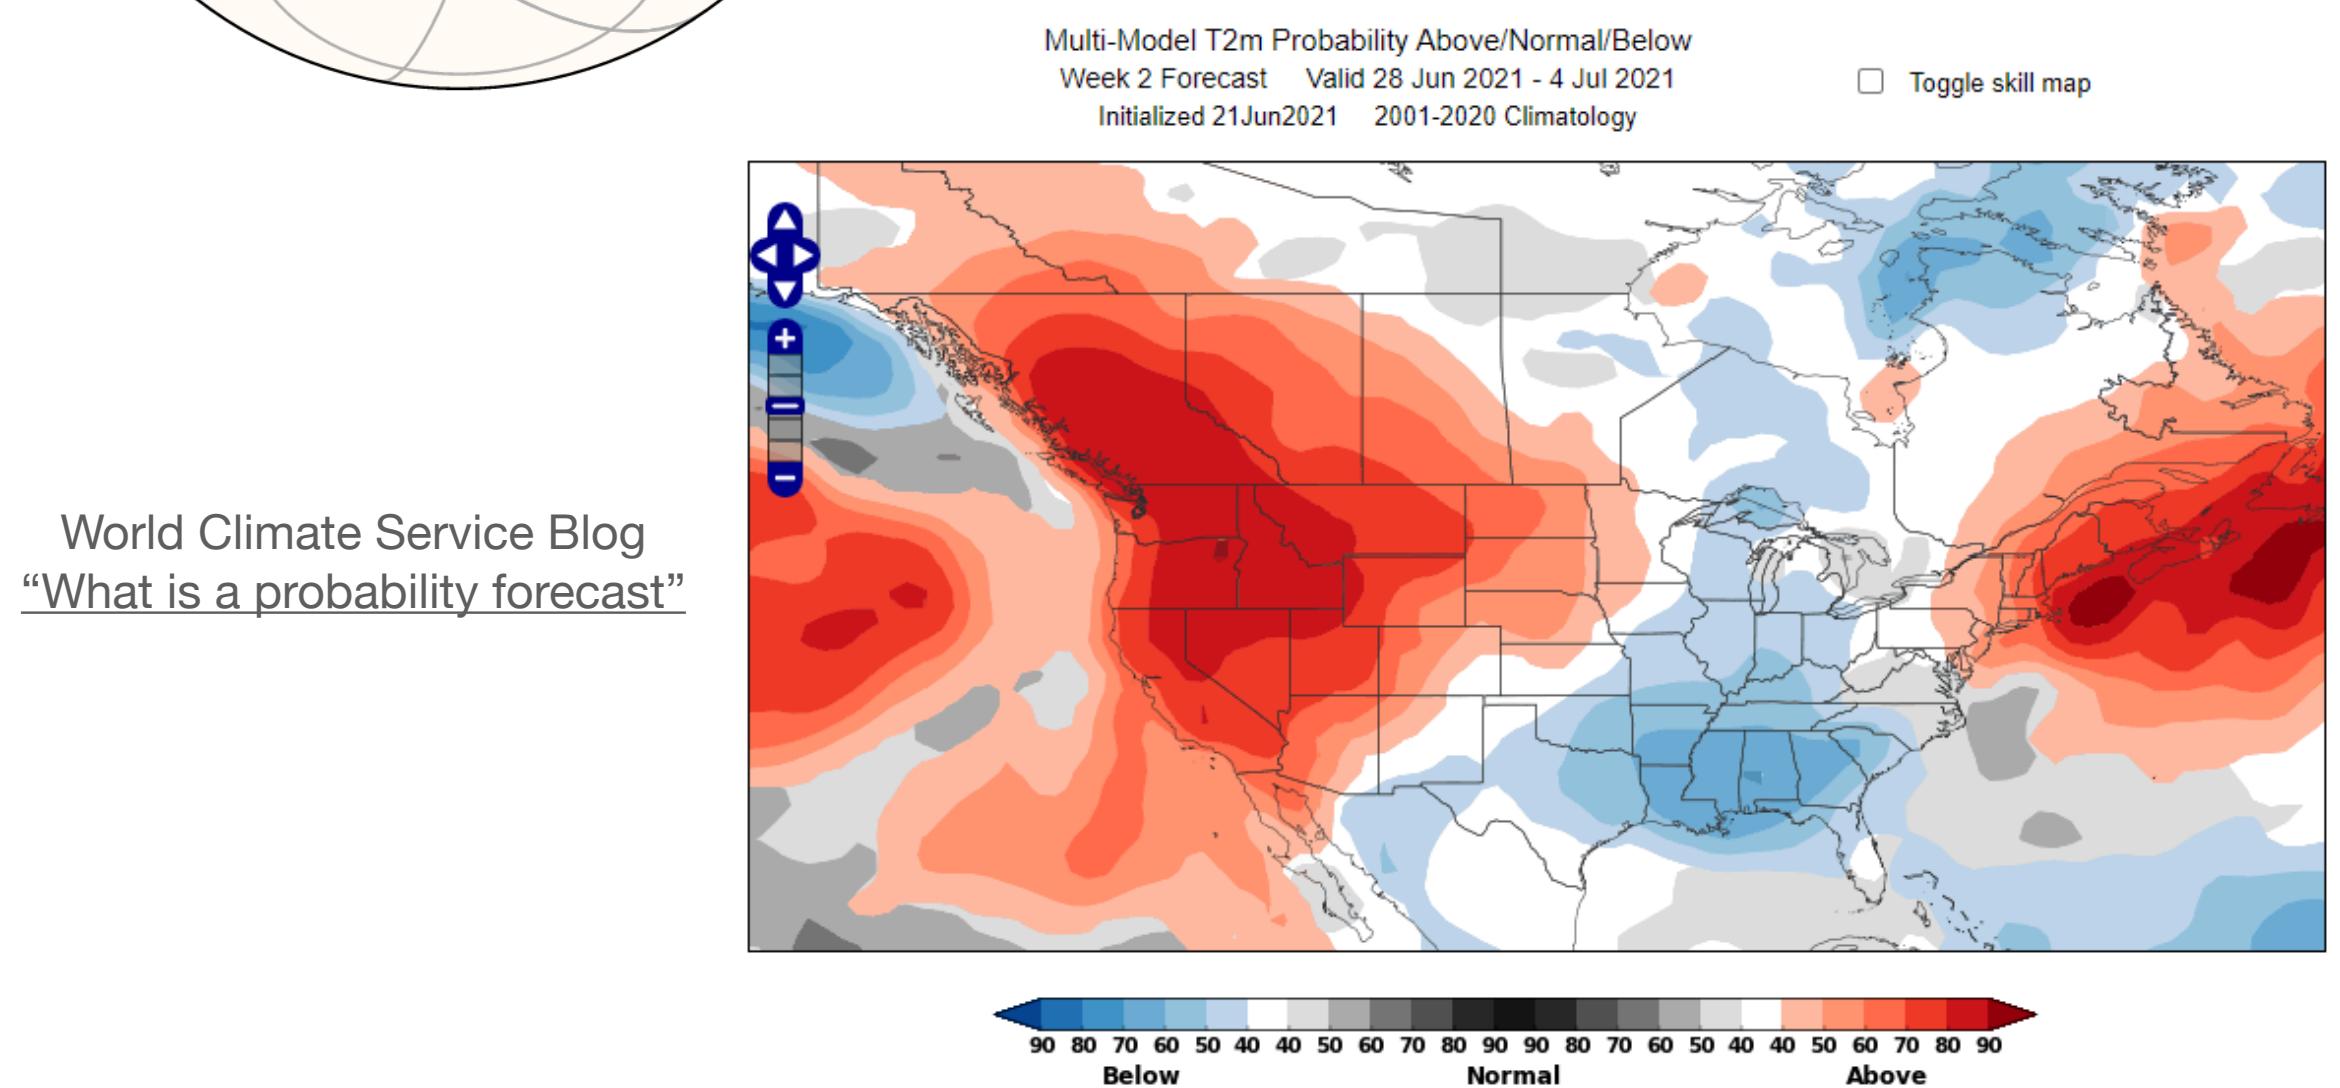


Why care about probability distributions?

- Data samples usually originate from some distribution
- Examples in natural sciences:
 - Positions of atoms in molecules
 - Parameters of gravitational waves
 - Temperature map
 - ...
- But: Distributions might be unknown, e.g. distribution over pixels in image



Midgley, Stimper, et al., Flow Annealed Importance Sampling Bootstrap, ICLR 2023



How can we learn a distribution over data?

- Train model to get

model distribution $p_{\text{model}} \approx \text{target distribution } p_{\text{target}}$

- What can we do with such models?

- Sample new data points $x \sim p_{\text{model}}(x)$



$$\sim p_{\text{dog-model}}(x)$$

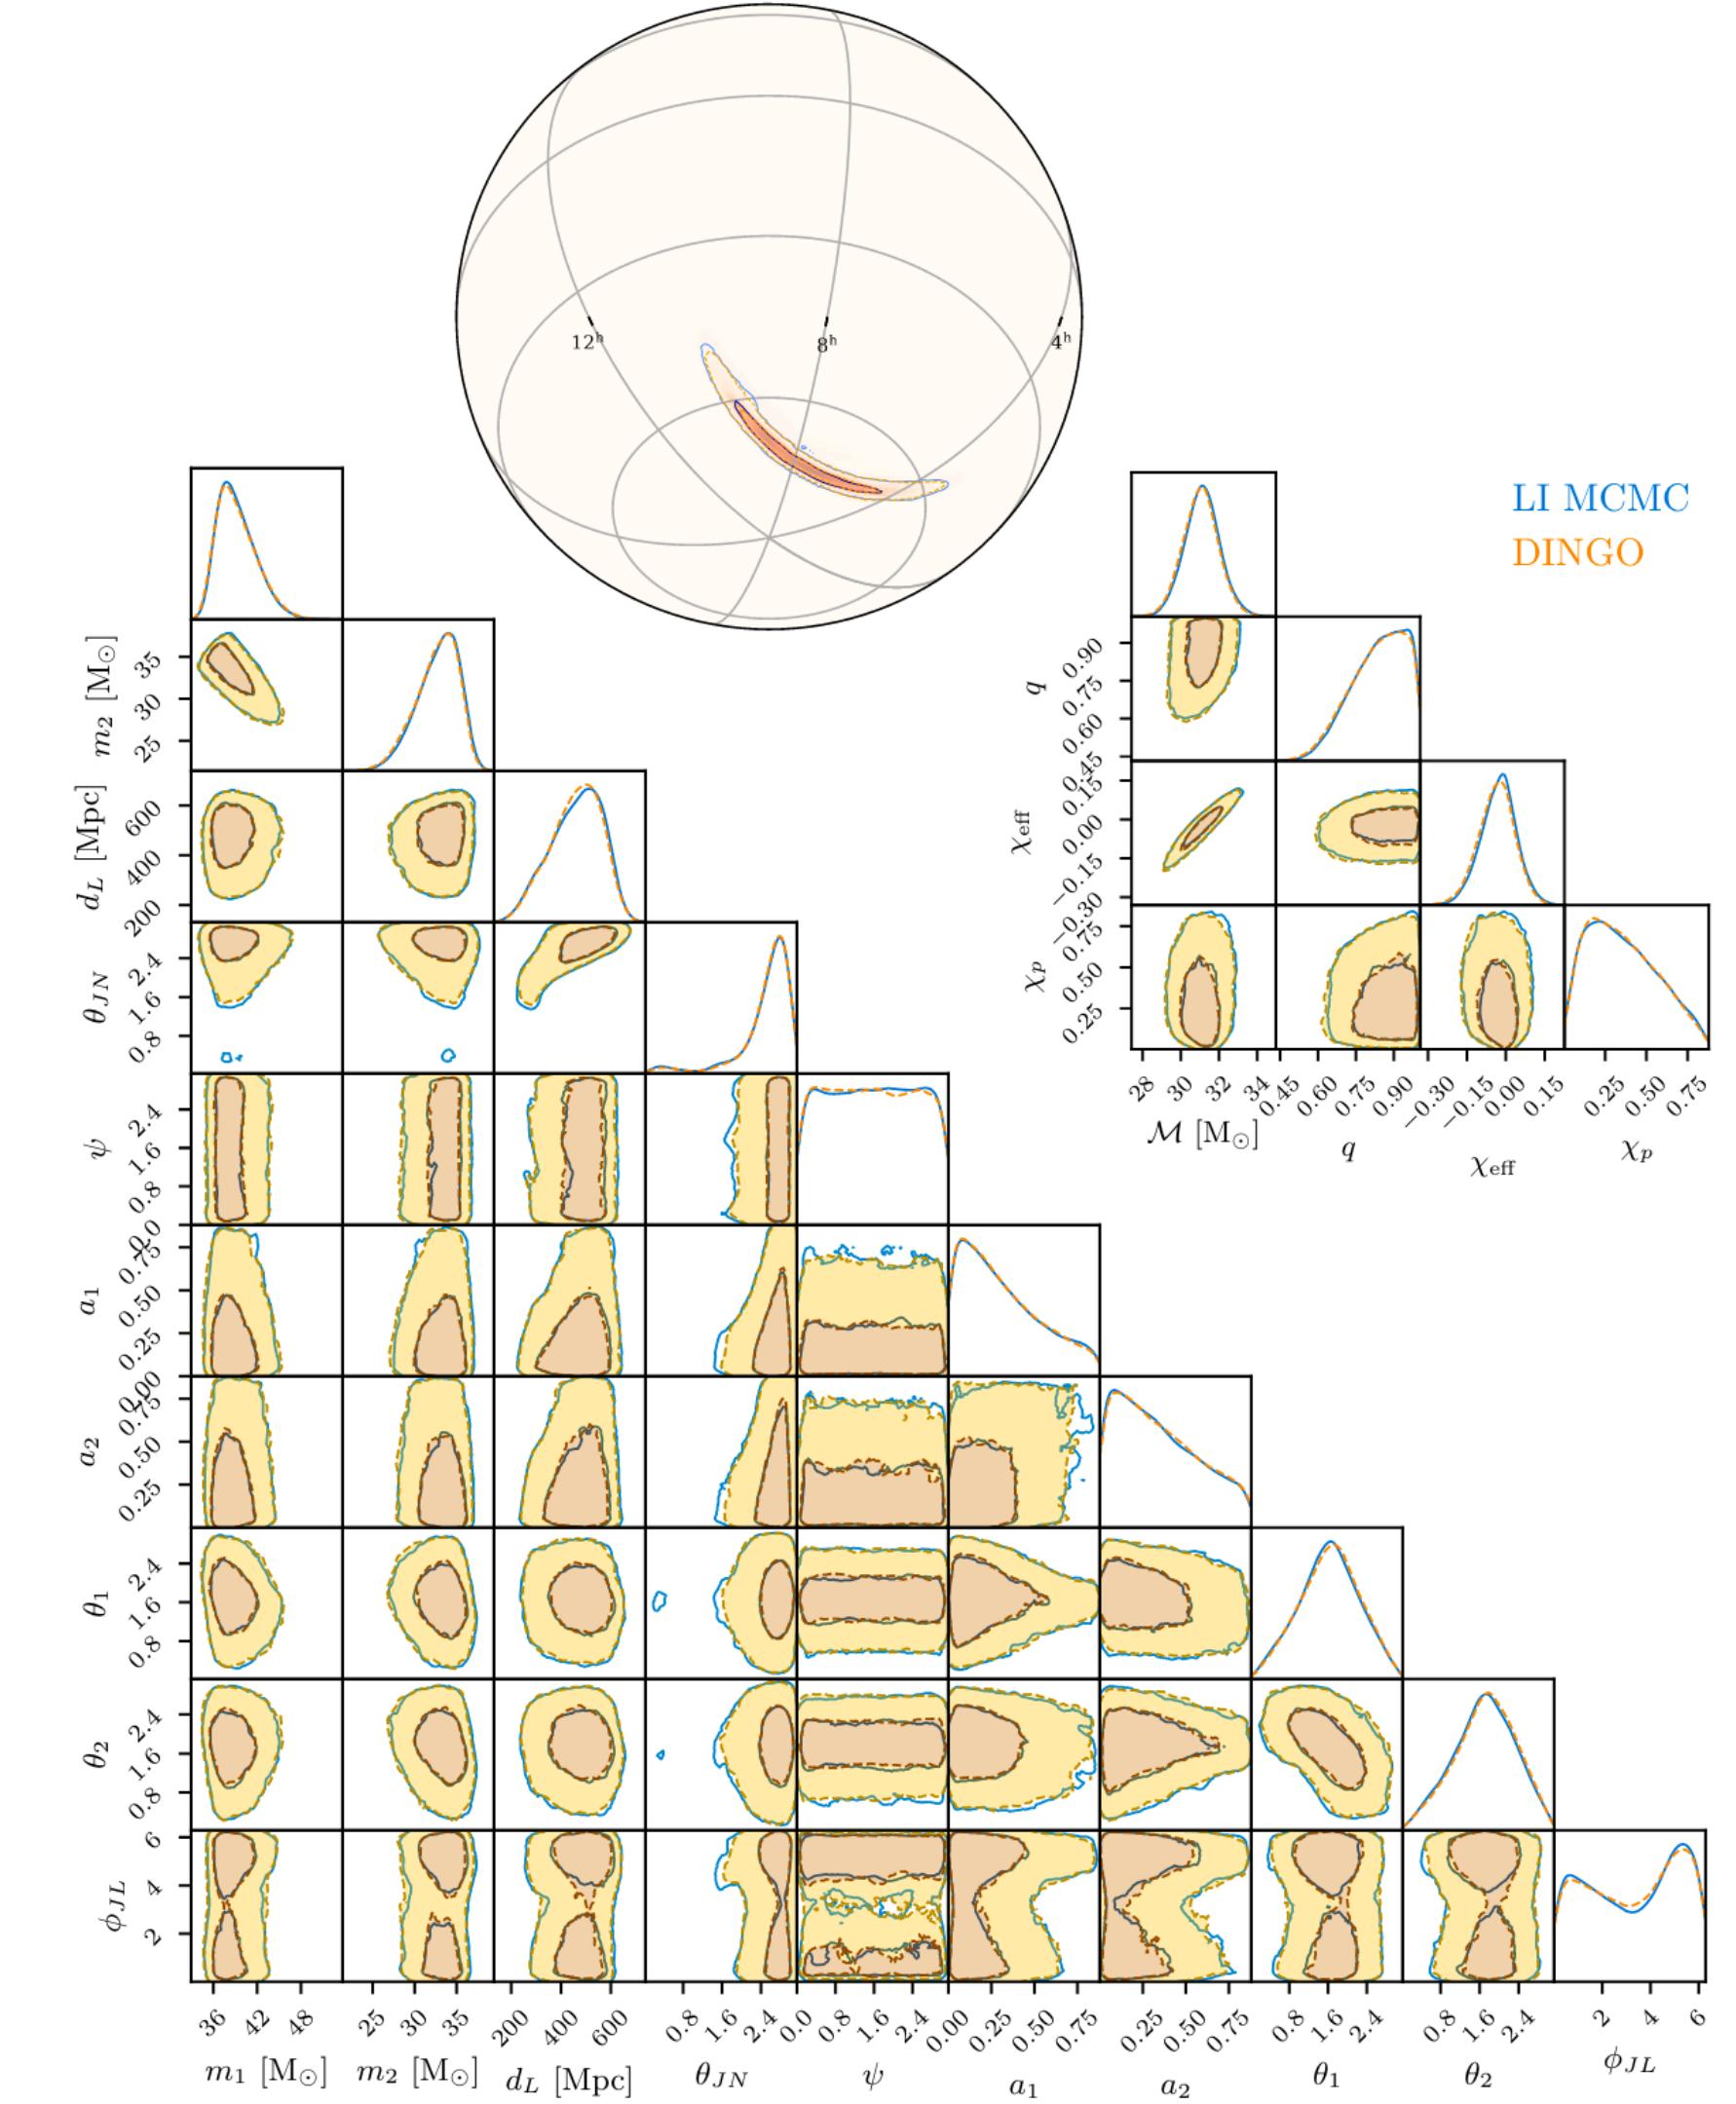
- Evaluation the model distribution $p_{\text{model}}(x_{\text{sample}})$ for a given data point x_{sample}

$$p_{\text{dog-model}} \left(\begin{array}{c} \text{small fluffy dog} \\ \text{standing on hind legs} \end{array} \right) = 0.87$$

Challenge: Complexity of Data Distribution

- Data distributions are usually extremely complex
 - Multi-modal
 - Unknown structure
 - High dimensional, e.g. pixels in image

→ powerful model required



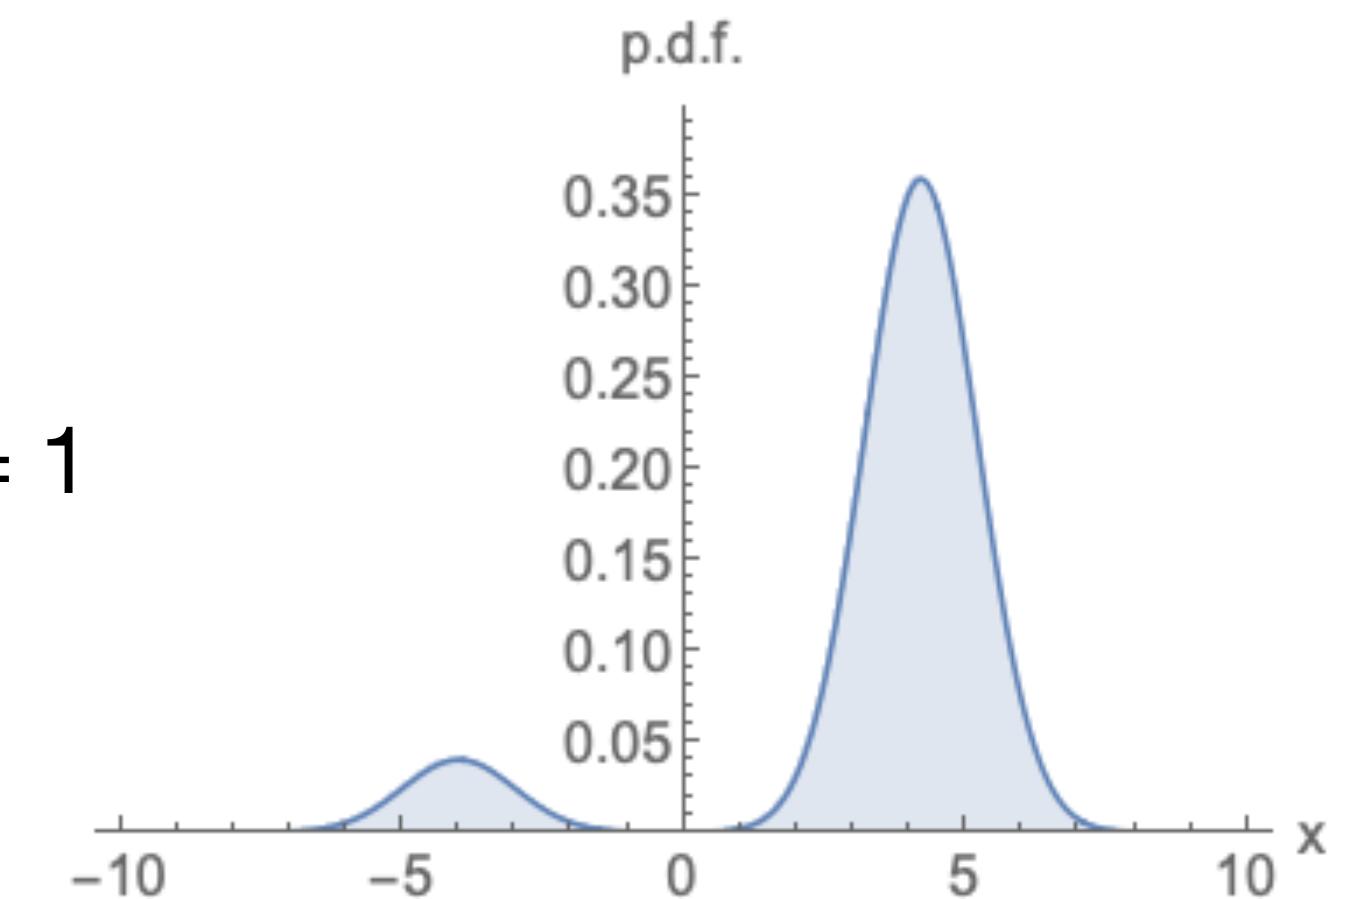
Dax, et al., Real-time gravitational-wave science
with neural posterior estimation, PRL2021

How to build a powerful model?

- Use a Deep Neural Network (DNN) f_θ
- Is the output $f_\theta(x)$ of the DNN for a sample x already a probability?
→ No, it does not fulfill the properties of a pdf.

1. Positive: $p(x) \geq 0 \forall x$

2. Normalized: $\int p(x) dx = 1$



How to build a powerful model?

- How can we adapt $f_\theta(x)$ such that it is a pdf?
 1. Positive $\rightarrow \exp(f_\theta(x))$
 2. Normalized $\rightarrow q_\theta(x) = \frac{\exp(f_\theta(x))}{Z_\theta}$
with normalization constant $Z_\theta = \int \exp(f_\theta(x))dx$
- Problem: calculating the normalization constant intractable for DNN

How to normalize? - Three Model Families

1. Approximate Z_θ

- Energy-based models

X Inaccurate probability

2. Use restricted NN architectures

- Variational autoencoders
- Autoregressive models
- Normalizing Flows

X Limited flexibility of the model

3. Model the generative process only

- GANs
- Diffusion models

X No/approximate evaluation of the probability

How to normalize? - Three Model Families

1. Approximate Z_θ

- Energy-based models

X Inaccurate probability

2. Use restricted NN architectures

- Variational autoencoders
- Autoregressive models
- **Normalizing Flows**

X Limited flexibility of the model

3. Model the generative process only

- GANs
- **Diffusion models**

X No/approximate evaluation of the probability

Normalizing Flows

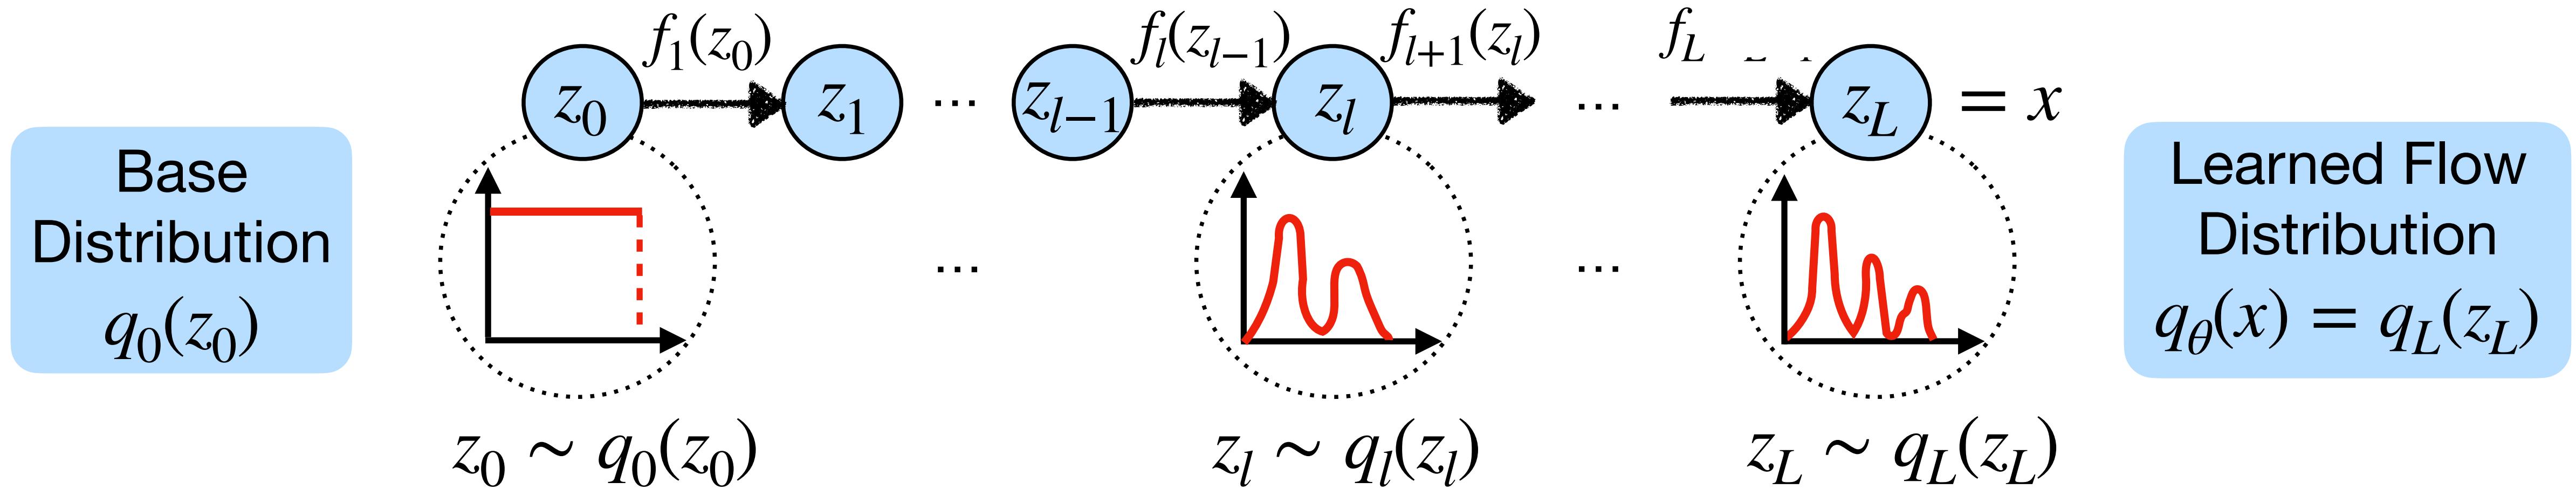
How to construct a Normalizing Flow?

Normalizing Flow



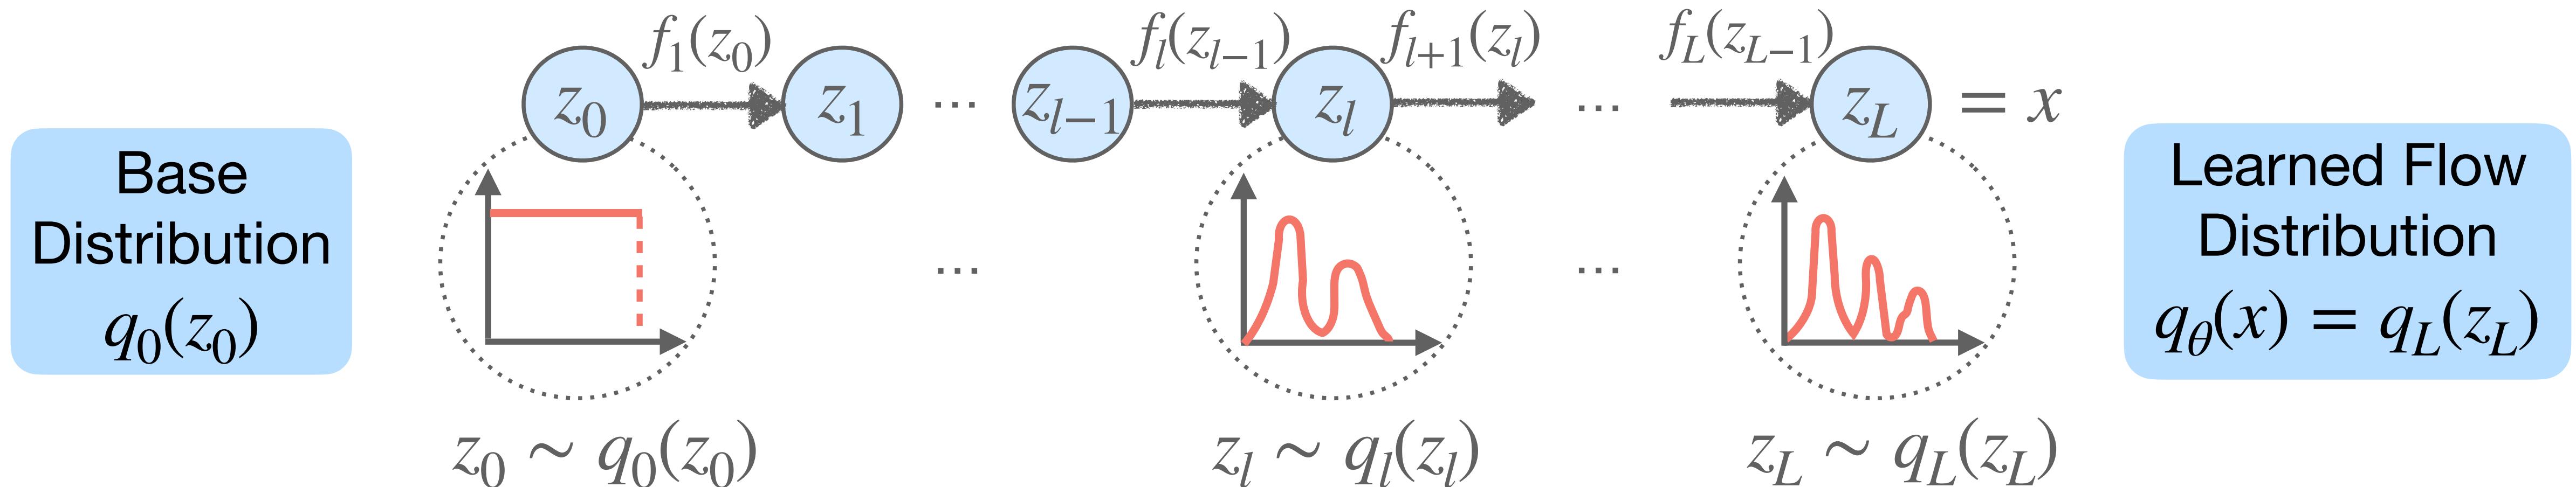
How to construct a Normalizing Flow?

Normalizing Flow = Sequence of invertible transformations:



How to sample from a Normalizing Flow?

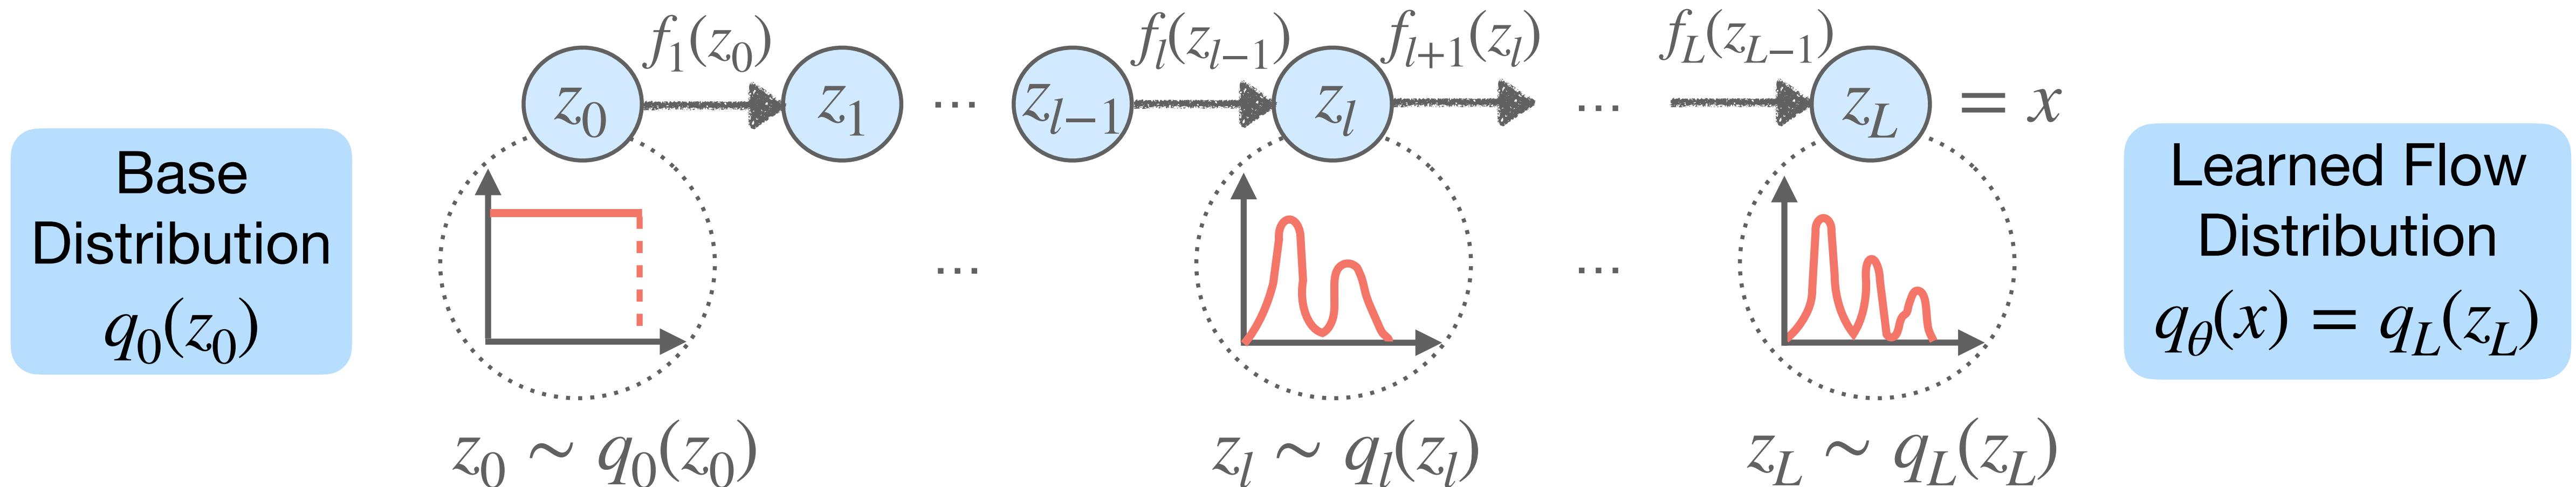
Normalizing Flow = Sequence of invertible transformations:



Sampling: $z_0 \sim q_0(z_0)$ $x = z_L = f_L \circ \dots \circ f_1(z_0)$

How to evaluate the density?

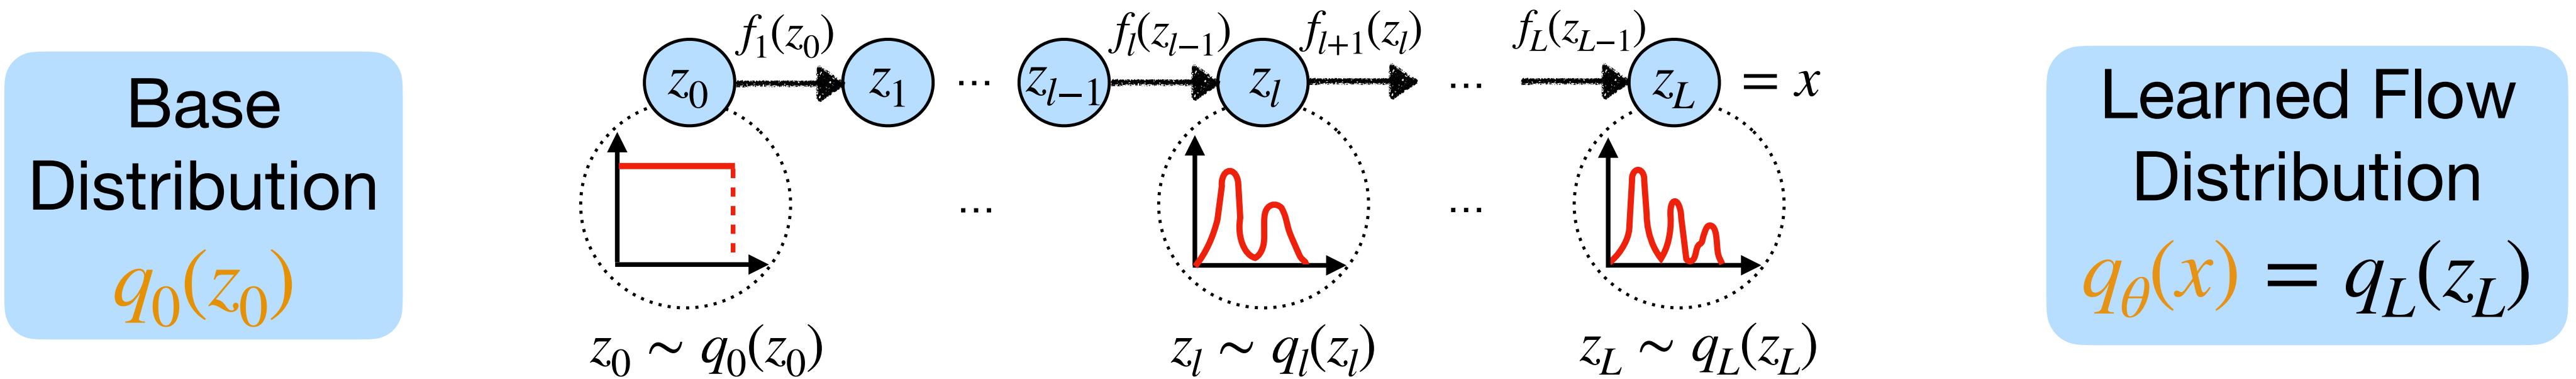
Normalizing Flow = Sequence of invertible transformations:



Evaluating the density: $z_0 = f_1^{-1} \circ \dots \circ f_L^{-1}(x)$ ← $x = z_L = f_L(z_{L-1})$

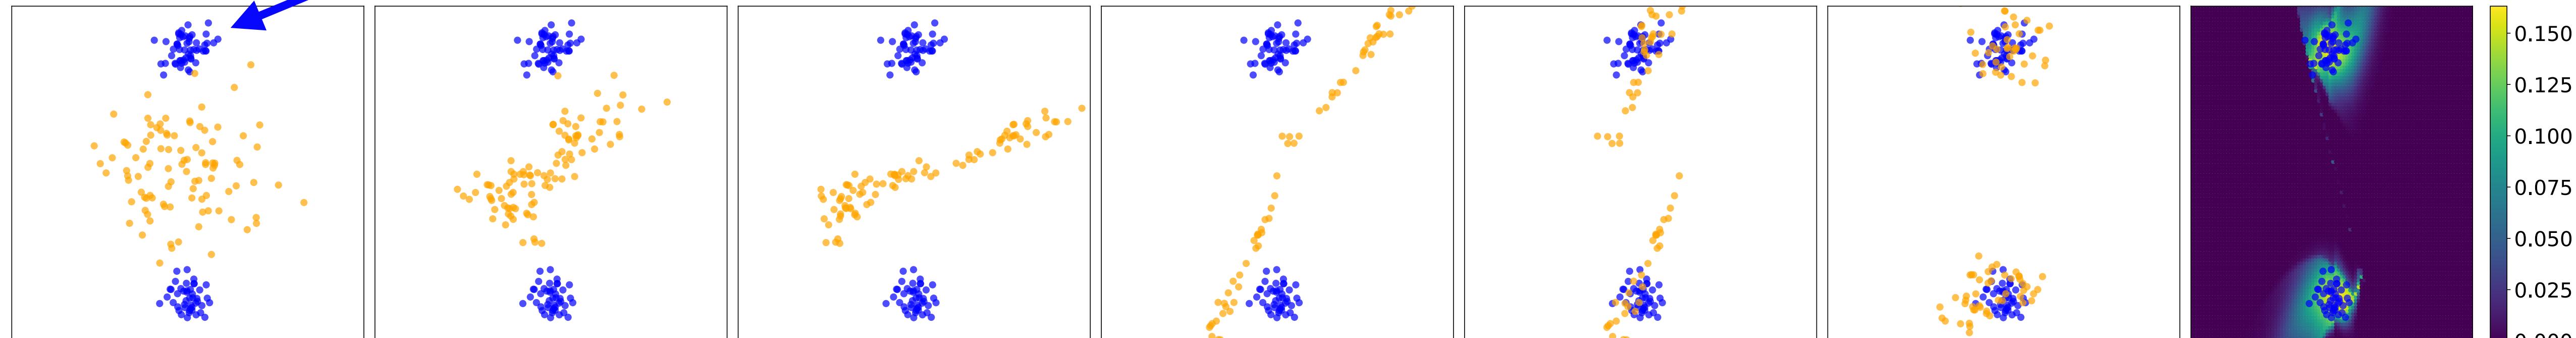
$$q(x) = q_0(f^{-1}(x)) \left| \det \left(\frac{\partial f^{-1}(x)}{\partial x} \right) \right| = q_0(f^{-1}(x)) \left| \det J_{f^{-1}} \right|$$

Example: Gaussian with two modes



Example:

Samples from target $p(x)$



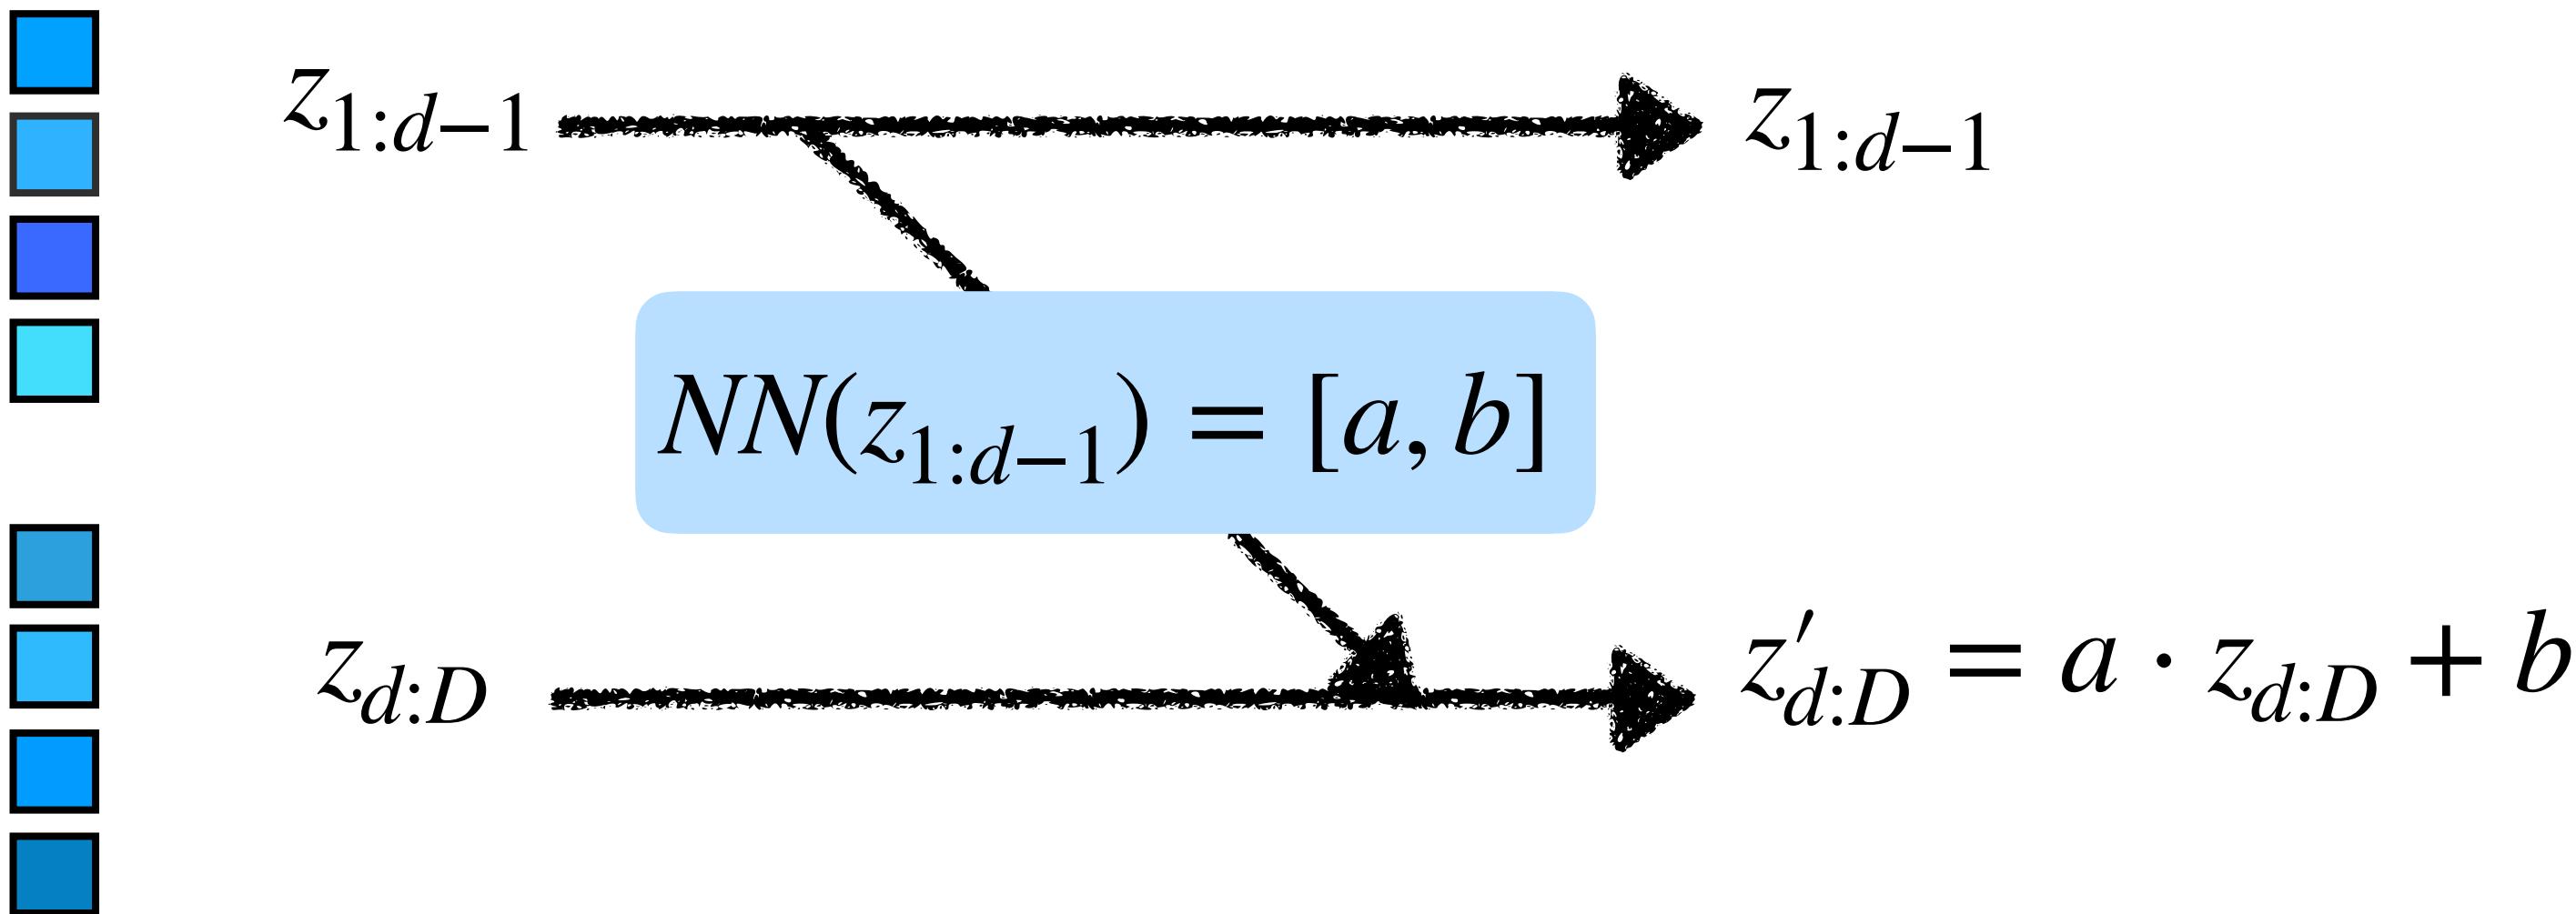
$z_0 \sim q_0(z_0)$

$z_5 \sim q_5(z_5)$ $q(x) = q_5(z_5)$

How to construct the transformations?

How to construct **invertible** transformations with **tractable** Jacobian:

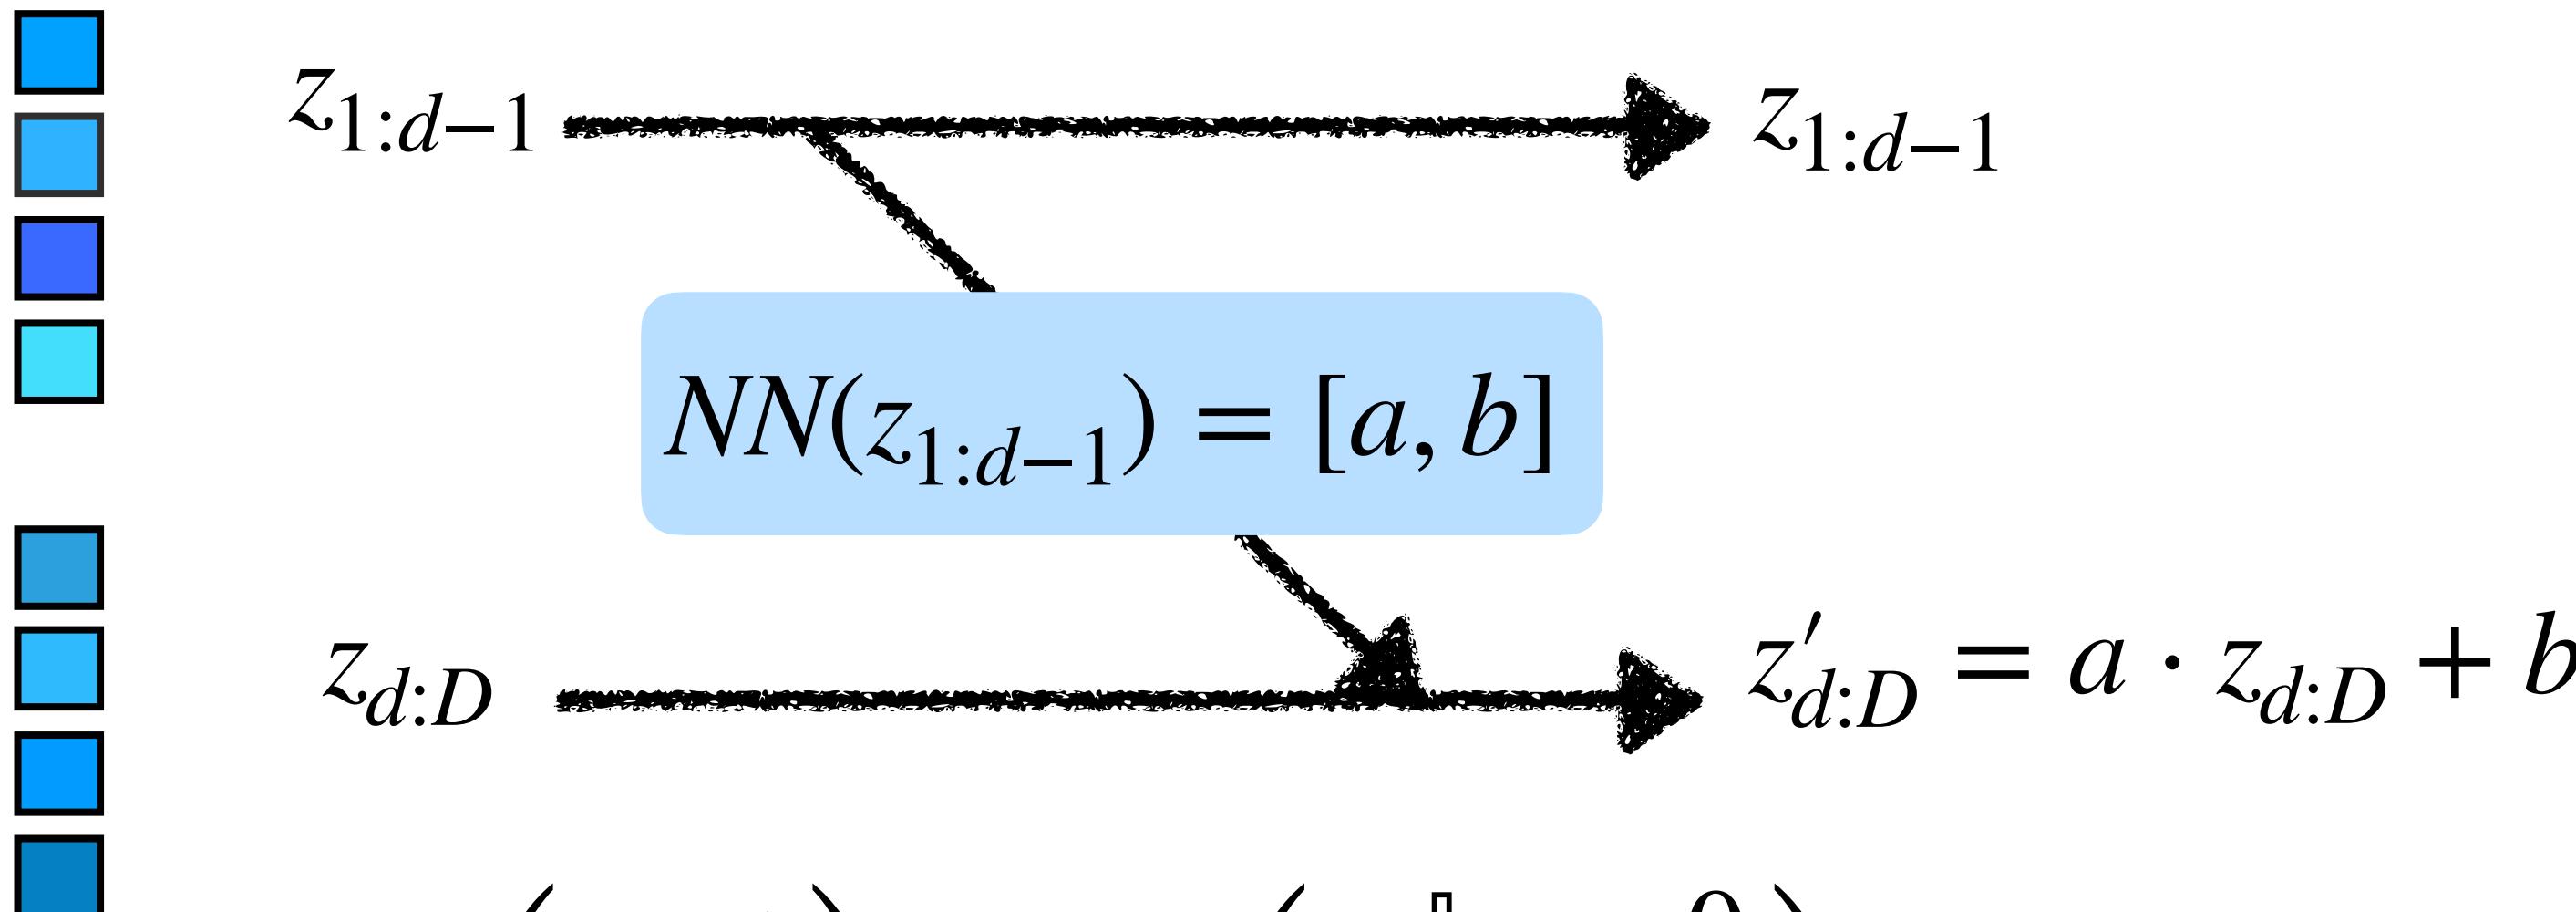
- Linear transformation: $z_i = a \cdot z_{i-1} + b \rightarrow$ learn a and b by NN
- Coupling layers: split $\vec{z} = (z_{1:d-1}, z_{d:D}) \in \mathbb{R}^D$



How to construct the transformations?

How to construct **invertible** transformations with **tractable** Jacobian:

- Linear transformation: $z_i = a \cdot z_{i-1} + b \rightarrow$ learn a and b by NN
- Coupling layers: split $\vec{z} = (z_{1:d-1}, z_{d:D}) \in \mathbb{R}^D$



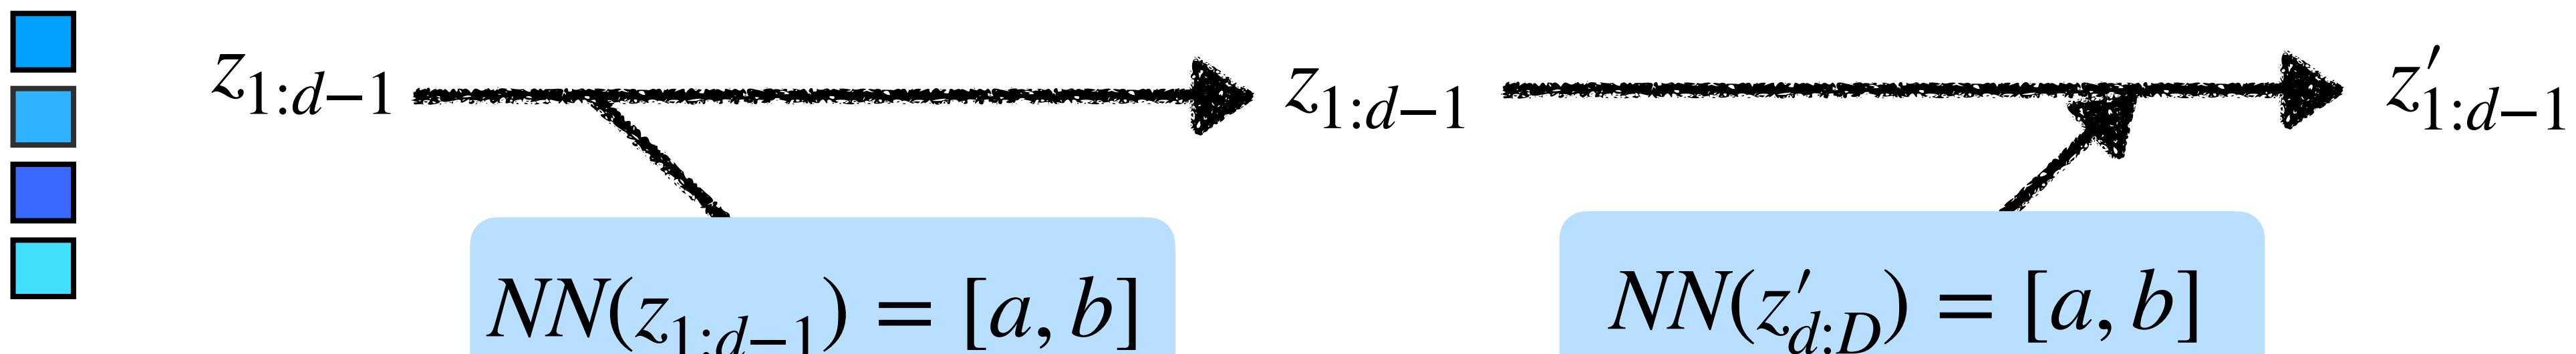
$$\det \left(\frac{\partial f^{-1}}{\partial z} \right) = \det \begin{pmatrix} \mathbb{I} & 0 \\ \frac{\partial z'_{d:D}}{\partial z_{1:d-1}} & \|a \end{pmatrix} = a^{D-d}$$

How to construct the transformations?

How to construct **invertible** transformations with **tractable** Jacobian:

- Linear transformation: $z_i = a \cdot z_{i-1} + b \rightarrow$ learn a and b by NN

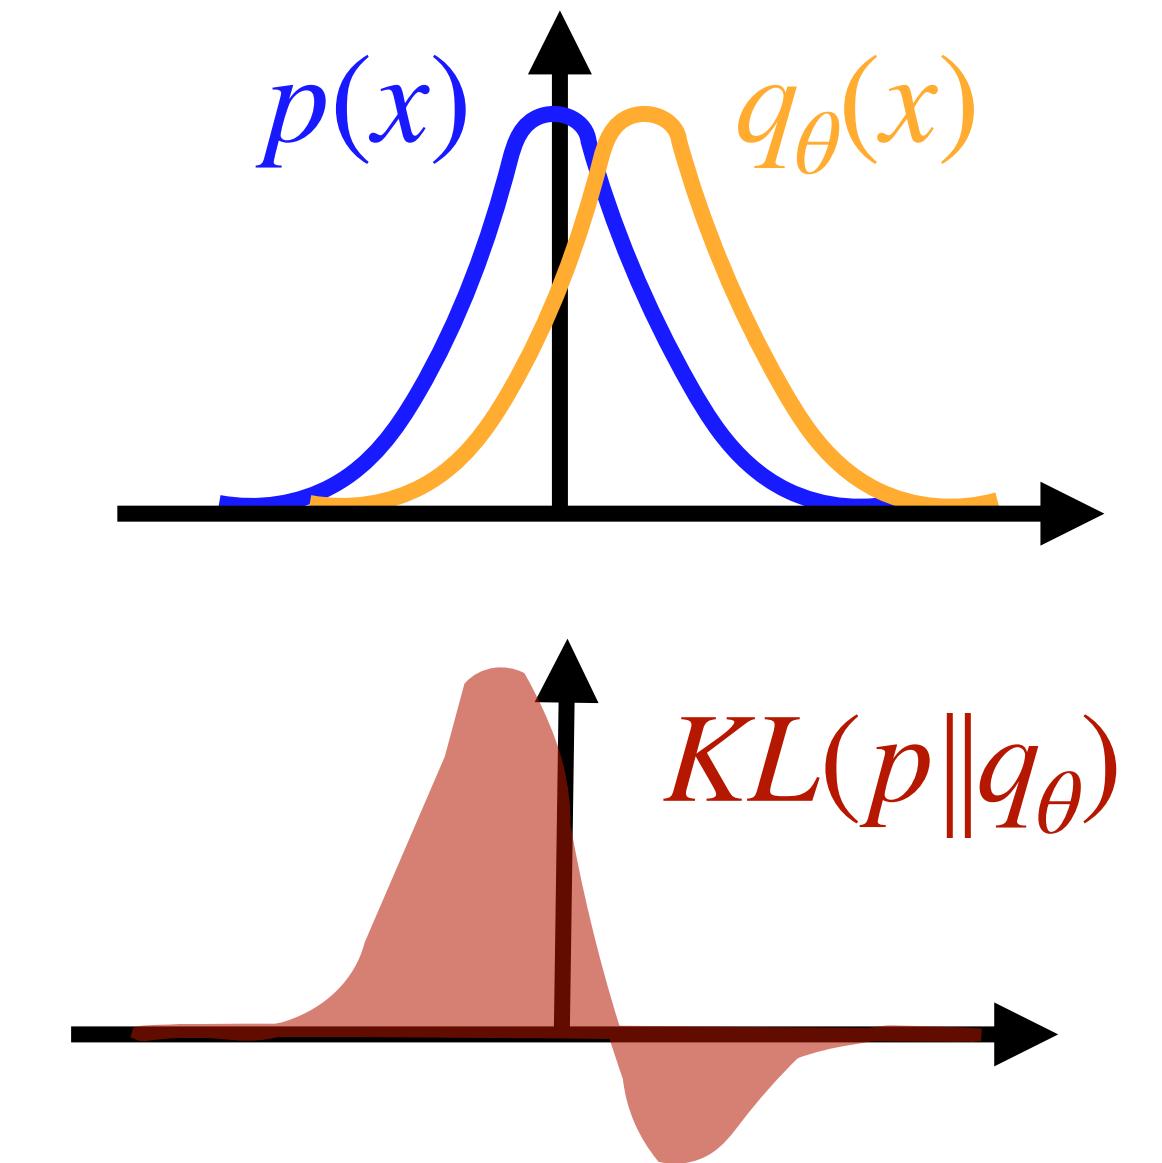
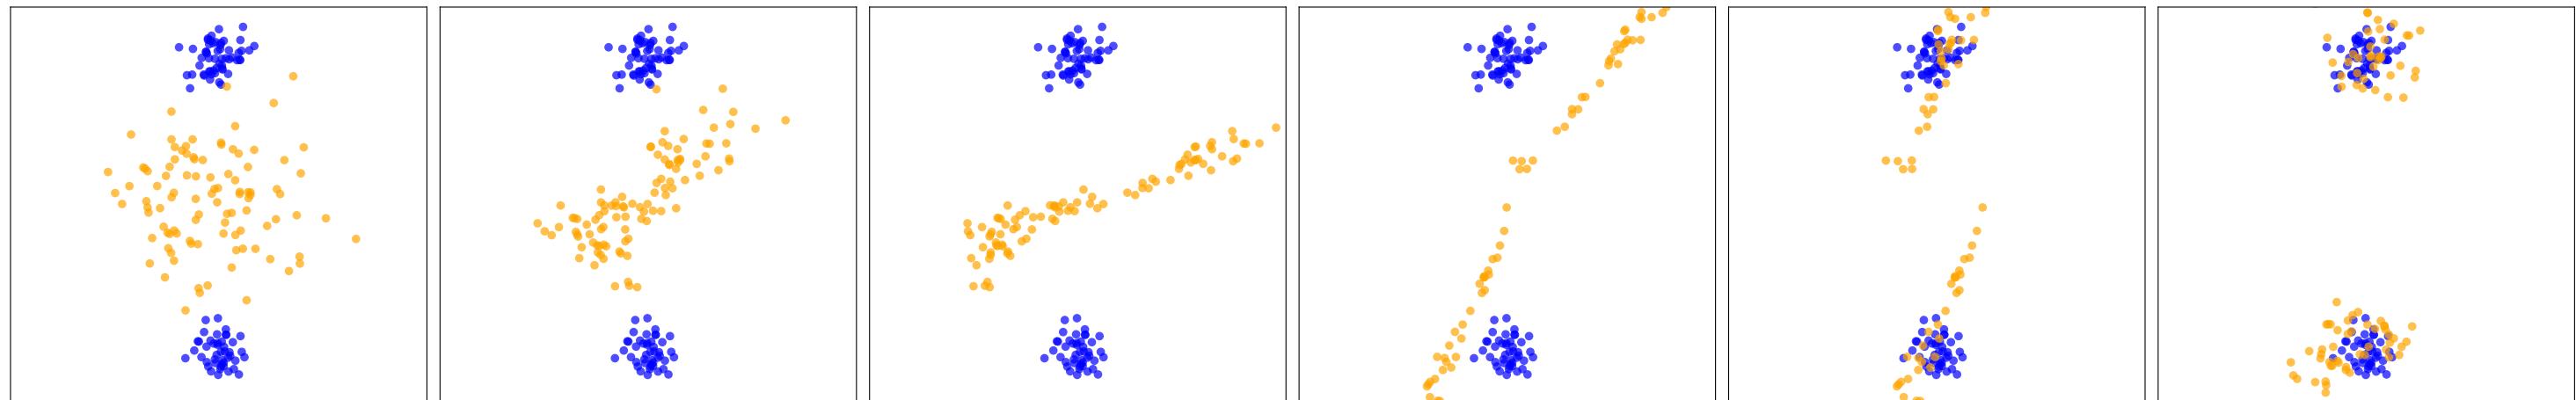
- Coupling layers: split $\vec{z} = (z_{1:d-1}, z_{d:D}) \in \mathbb{R}^D$



$$\det \left(\frac{\partial f^{-1}}{\partial z} \right) = \det \begin{pmatrix} \mathbb{I} & 0 \\ \frac{\partial z'_{d:D}}{\partial z_{1:d-1}} & \|a \end{pmatrix} = a^{D-d}$$

How to train a normalizing flow?

How can we compare two distributions?



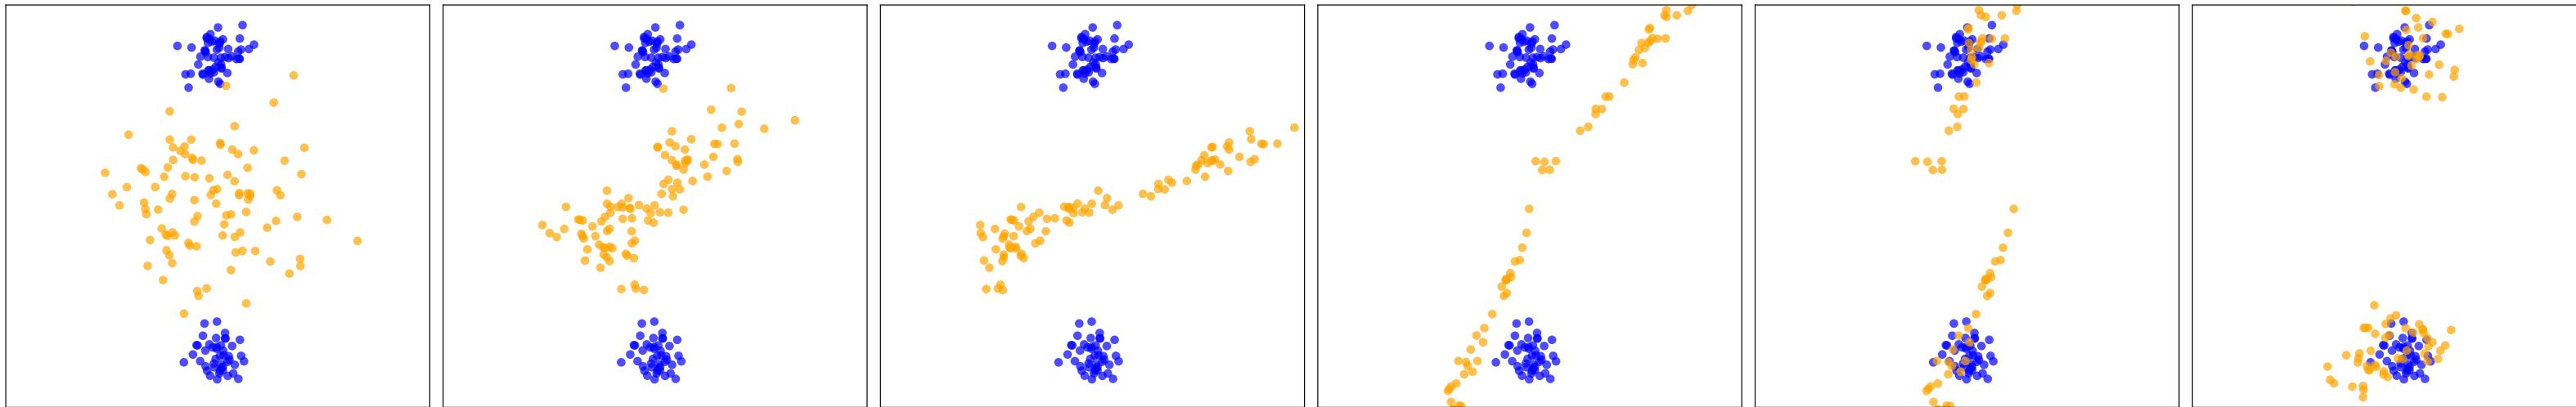
Kullback-Leibler divergence

$$\mathcal{L} = KL(p \parallel q_\theta) = \mathbb{E}_{x \sim p(x)} \left[\log \frac{p(x)}{q_\theta(x)} \right] = - \sum_{i=1}^N \log q_\theta(x_i) + \text{const.}$$

→ minimize negative log-likelihood by evaluating the density of training samples

Normalizing flows: 3 Take-Aways

1. Normalizing flow = sequence of invertible transformations



2. Transformation constructed to have tractable Jacobian
→ Exact density evaluation, but limited architecture
3. Trained by evaluating and minimizing the negative log-likelihood of training data samples

How to normalize? - Three Model Families

1. Approximate Z_θ

- Energy-based models

X Inaccurate probability

2. Use restricted NN architectures

- Variational autoencoders
- Autoregressive models
- **Normalizing Flows** ✓

X Limited flexibility of the model

3. Model the generative process only

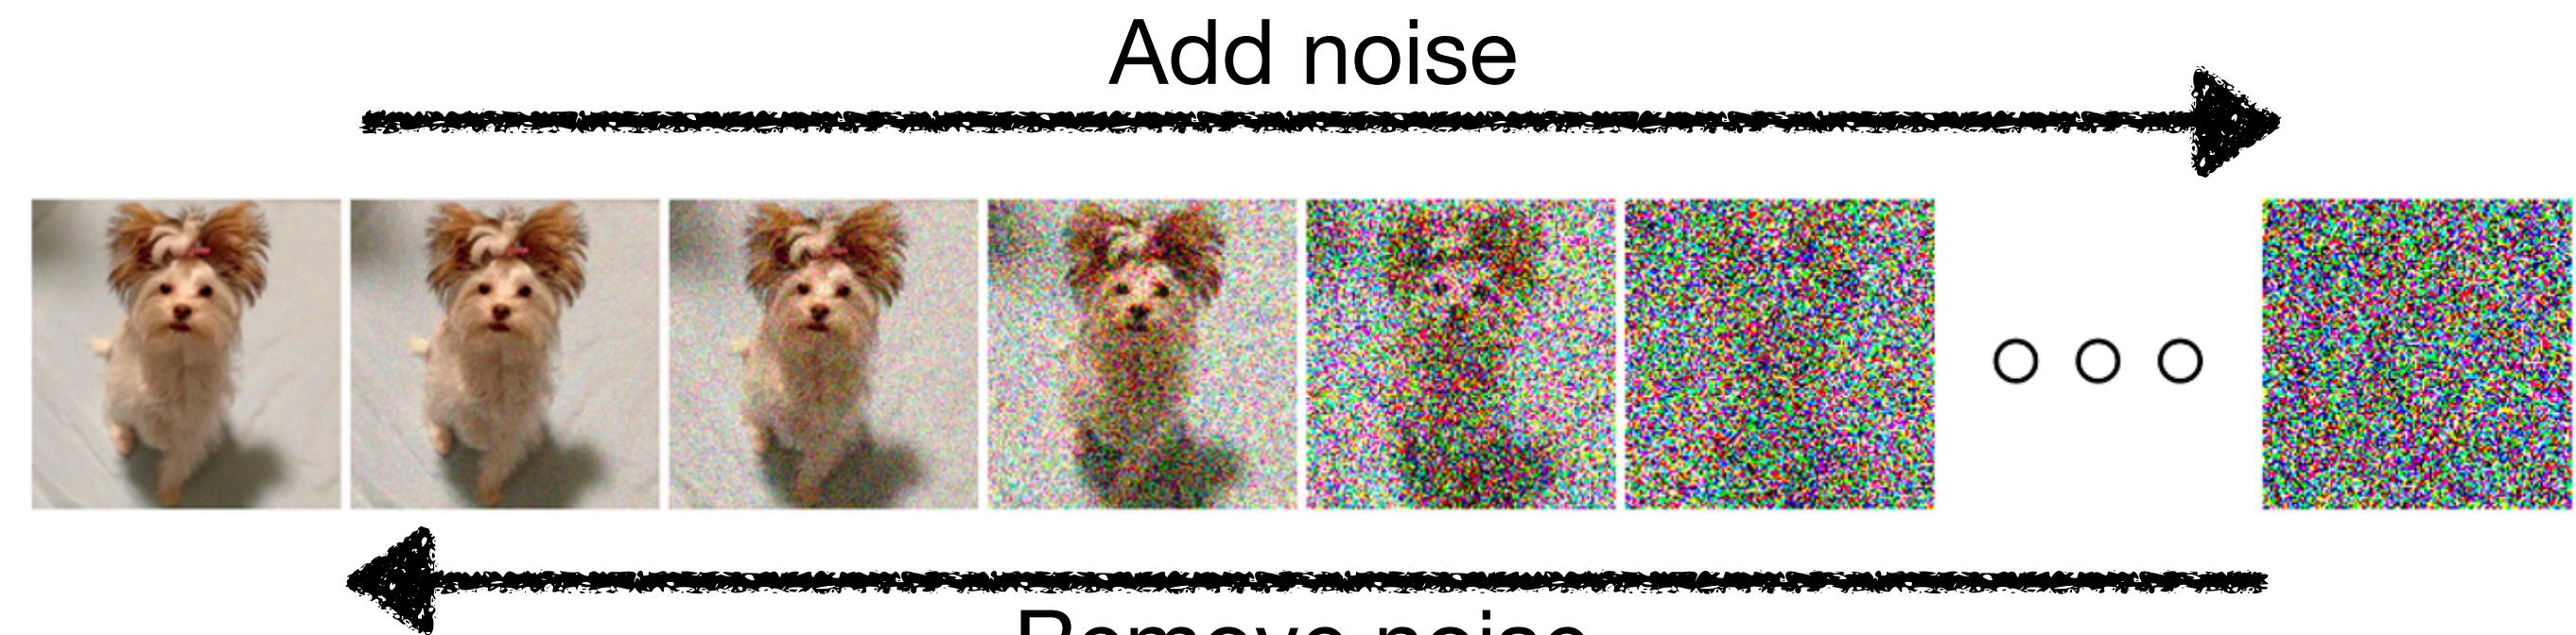
- GANs
- **Diffusion models**

X No/approximate evaluation of the probability

Diffusion Models

What is a diffusion model?

- Characterized by
 - progressively adding noise to the data
 - learning the reverse process



- Members of the diffusion family:
 - Denoising Diffusion Probabilistic Model
 - Score-Based Generative Models
 - Stochastic Differential Equations

What is a diffusion model?

- Characterized by
 - progressively adding noise to the data
 - learning the reverse process



- Members of the diffusion family:
 - **Denoising Diffusion Probabilistic Model**
 - Score-Based Generative Models
 - Stochastic Differential Equations

How does a diffusion model work?

- Forward process: add Gaussian noise to data



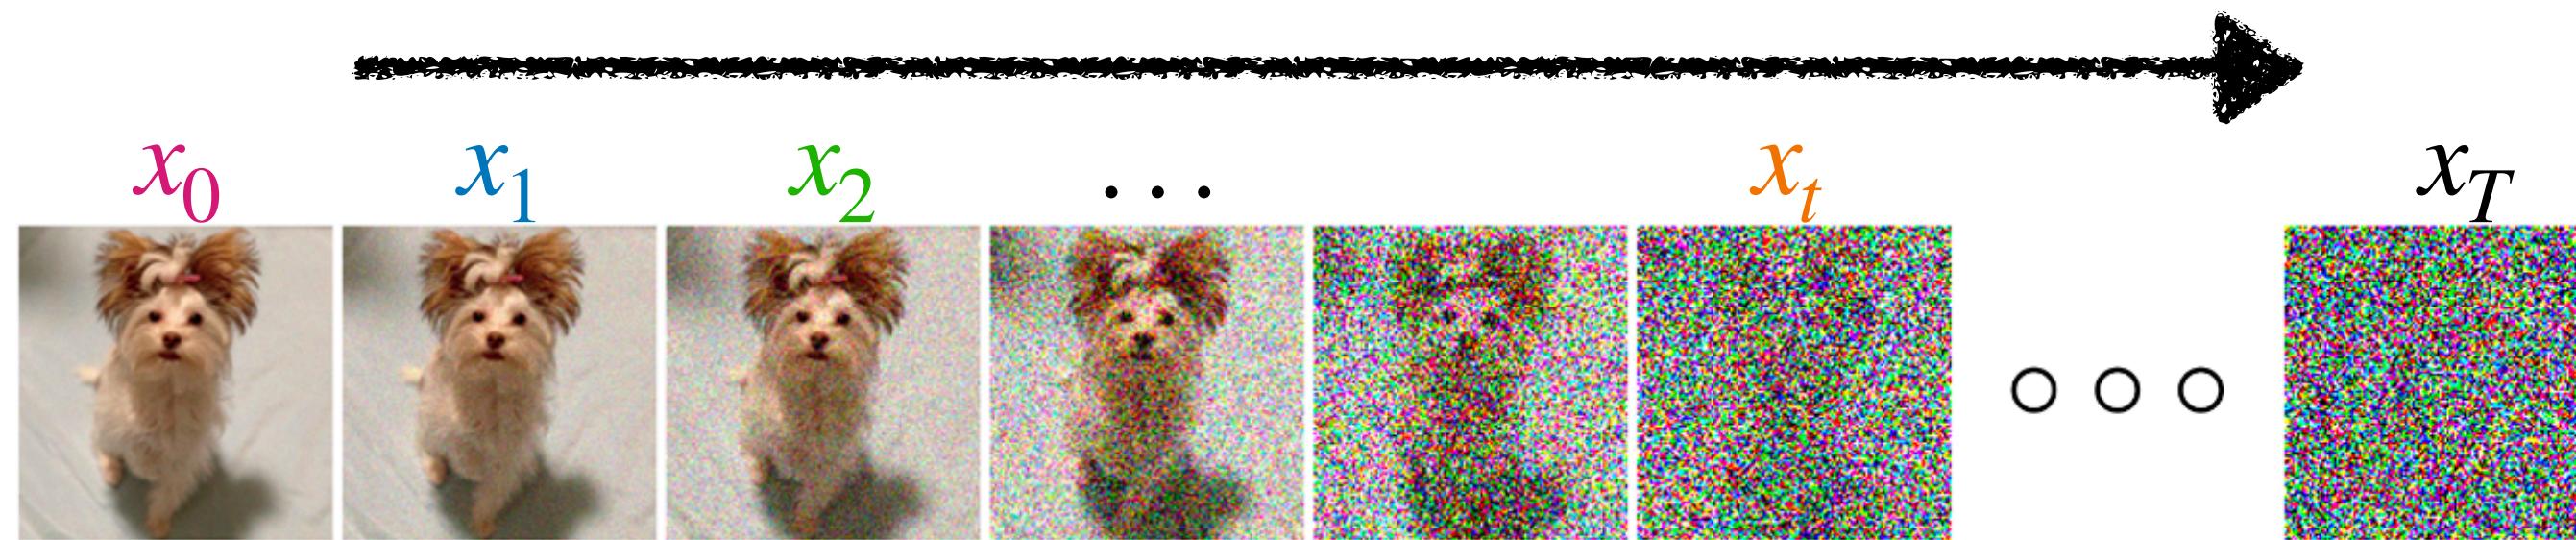
$$\begin{aligned} p(x_0) \\ \xrightarrow{} p(x_1|x_0) &= \mathcal{N}(x_1 | \sqrt{\alpha_1}x_0, (1 - \alpha_1)\mathbb{I}) \\ \xrightarrow{} p(x_2|x_1, x_0) &= \mathcal{N}(x_2 | \sqrt{\alpha_2}x_1, (1 - \alpha_2)\mathbb{I}) \\ &\vdots \end{aligned}$$

- Why Gaussians?

$$p(x_t|x_0) = \mathcal{N}\left(x_t \mid \sqrt{\prod_{s=0}^t \alpha_s} x_0, \left(1 - \prod_{s=0}^t \alpha_s\right)\mathbb{I}\right) \equiv \mathcal{N}\left(x_t \mid \sqrt{\bar{\alpha}_t} x_0, (1 - \sqrt{\bar{\alpha}_t})\mathbb{I}\right)$$

How does a diffusion model work?

- Forward process: add Gaussian noise to data



- Why Gaussians?
→ reparametrization trick: sample noise $\epsilon \sim \mathcal{N}(\epsilon | 0, \mathbb{I})$

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + (1 - \sqrt{\bar{\alpha}_t}) \epsilon$$

How does a diffusion model work?

- Reverse process: “Remove” noise → learn by DNN



- How to remove noise?

$$x_t \rightarrow NN(x_t, t) = [\mu_\theta, \Sigma_\theta]$$



$$q_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1} | \mu_\theta(x_t, t), \Sigma_\theta(x_t, t)) \rightarrow x_{t-1} \sim q_\theta(x_{t-1}|x_t)$$

How to sample from a diffusion model?

- Sample random noise



$$x_T \sim \mathcal{N}(x_T | 0, \mathbb{I})$$

- Go through reverse process until $t = 0$:

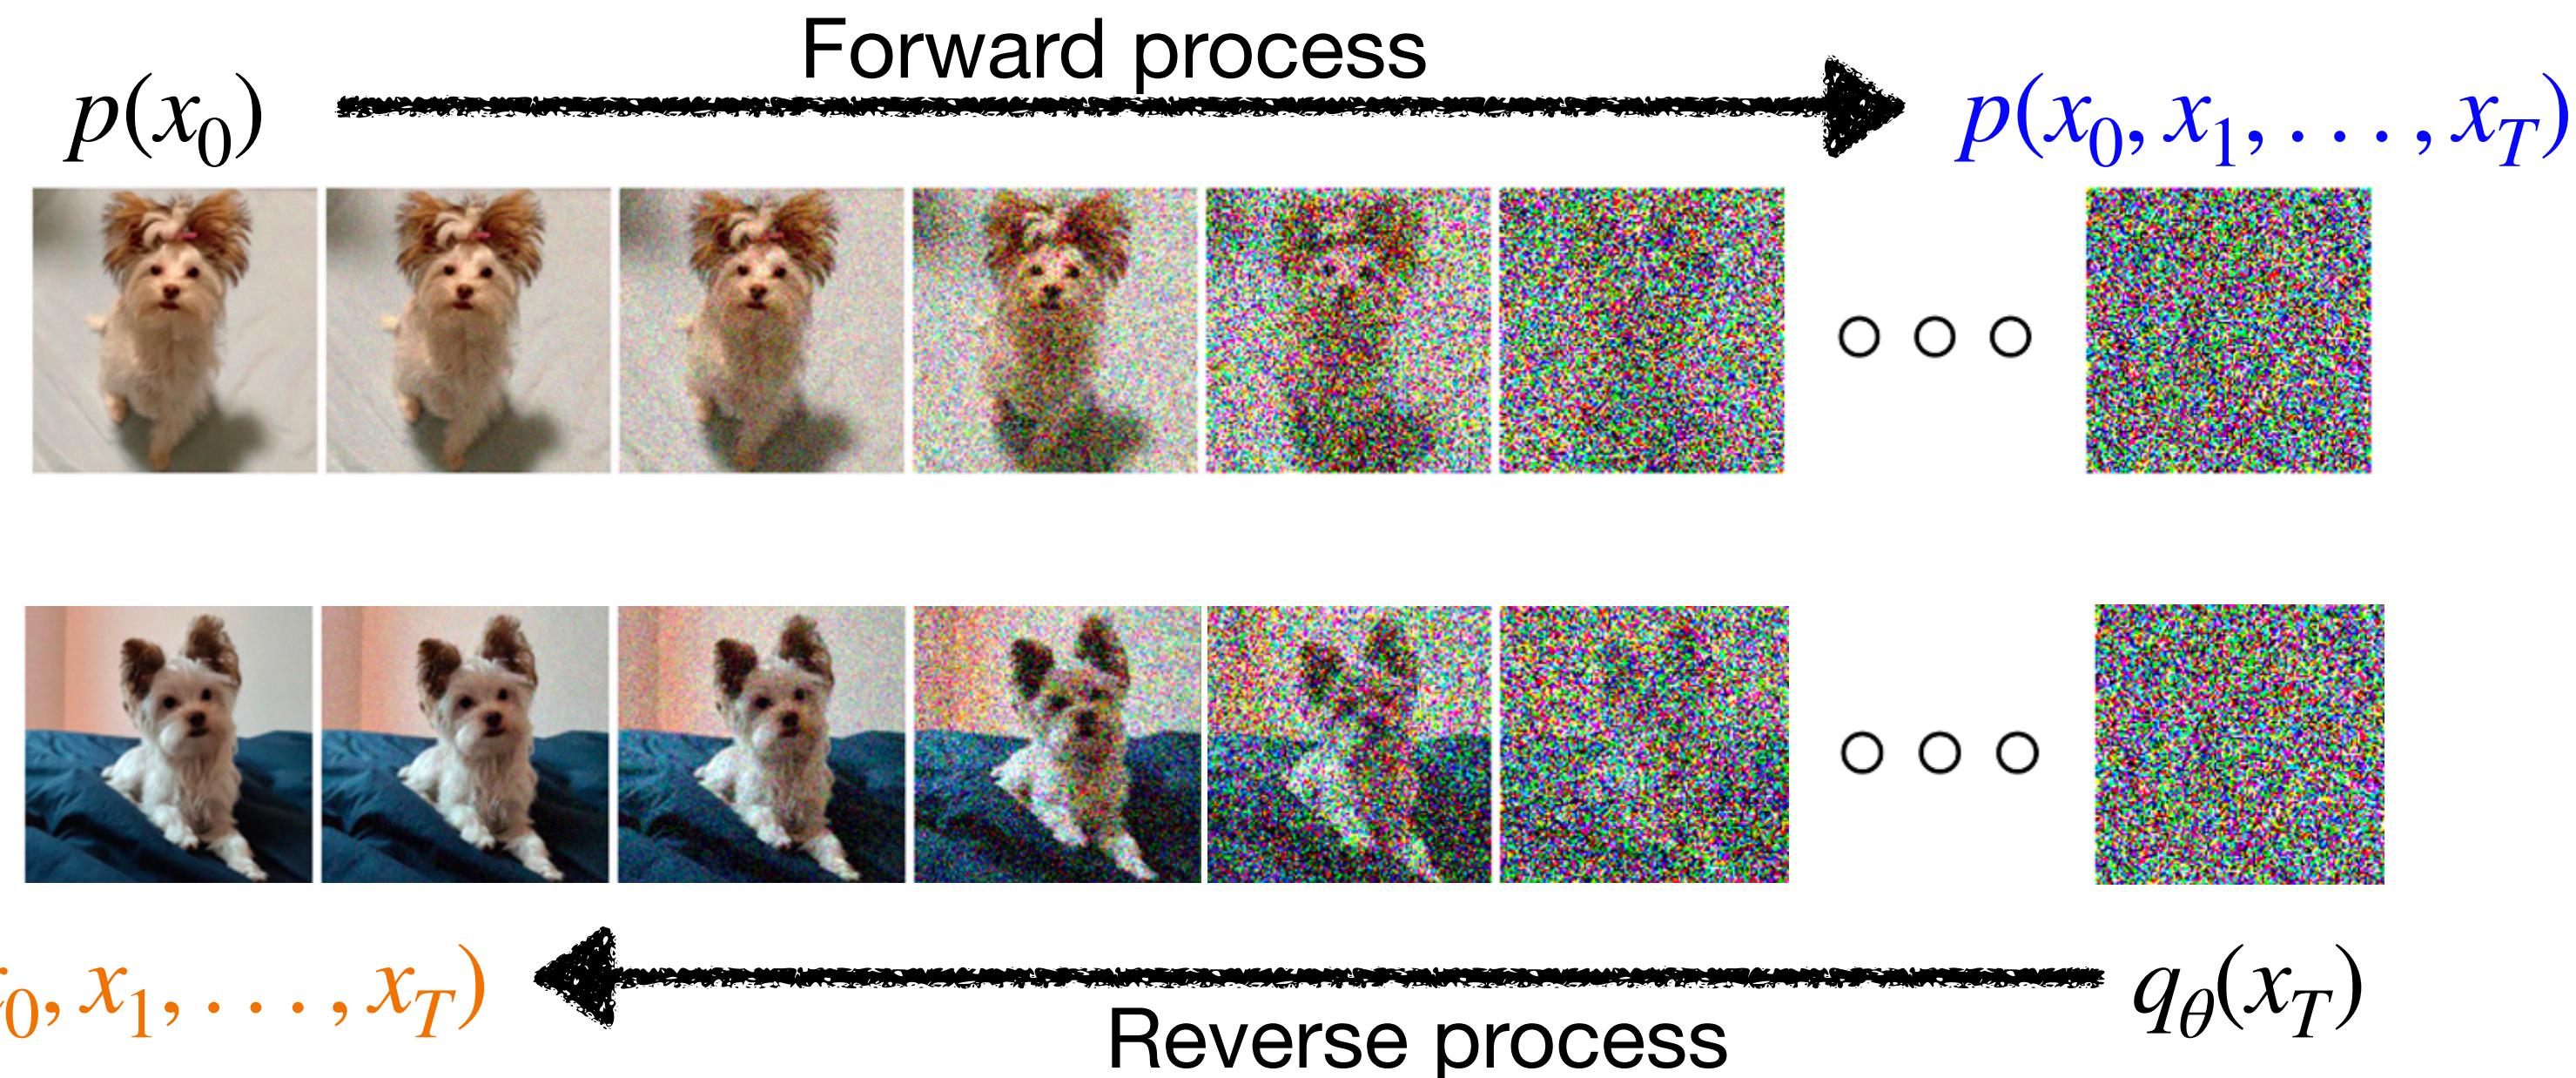
$$x_1 \rightarrow NN(x_1, t) = [\mu_\theta, \Sigma_\theta]$$



$$q_\theta(x_0|x_1) = \mathcal{N}(x_0 | \mu_\theta(x_1, t), \Sigma_\theta(x_1, t)) \rightarrow x_0 \sim q_\theta(x_0|x_1)$$

How to train a diffusion model?

- Compare joint distributions with Kullback-Leibler divergence:



$$\mathcal{L} = KL(p(x_0, x_1, \dots, x_T) \| q_{\theta}(x_0, \dots, x_T)) = - \mathbb{E}_{p(x_0, x_1, \dots, x_T)} [\log q_{\theta}(x_0, x_1, \dots, x_T)] + \text{const.}$$

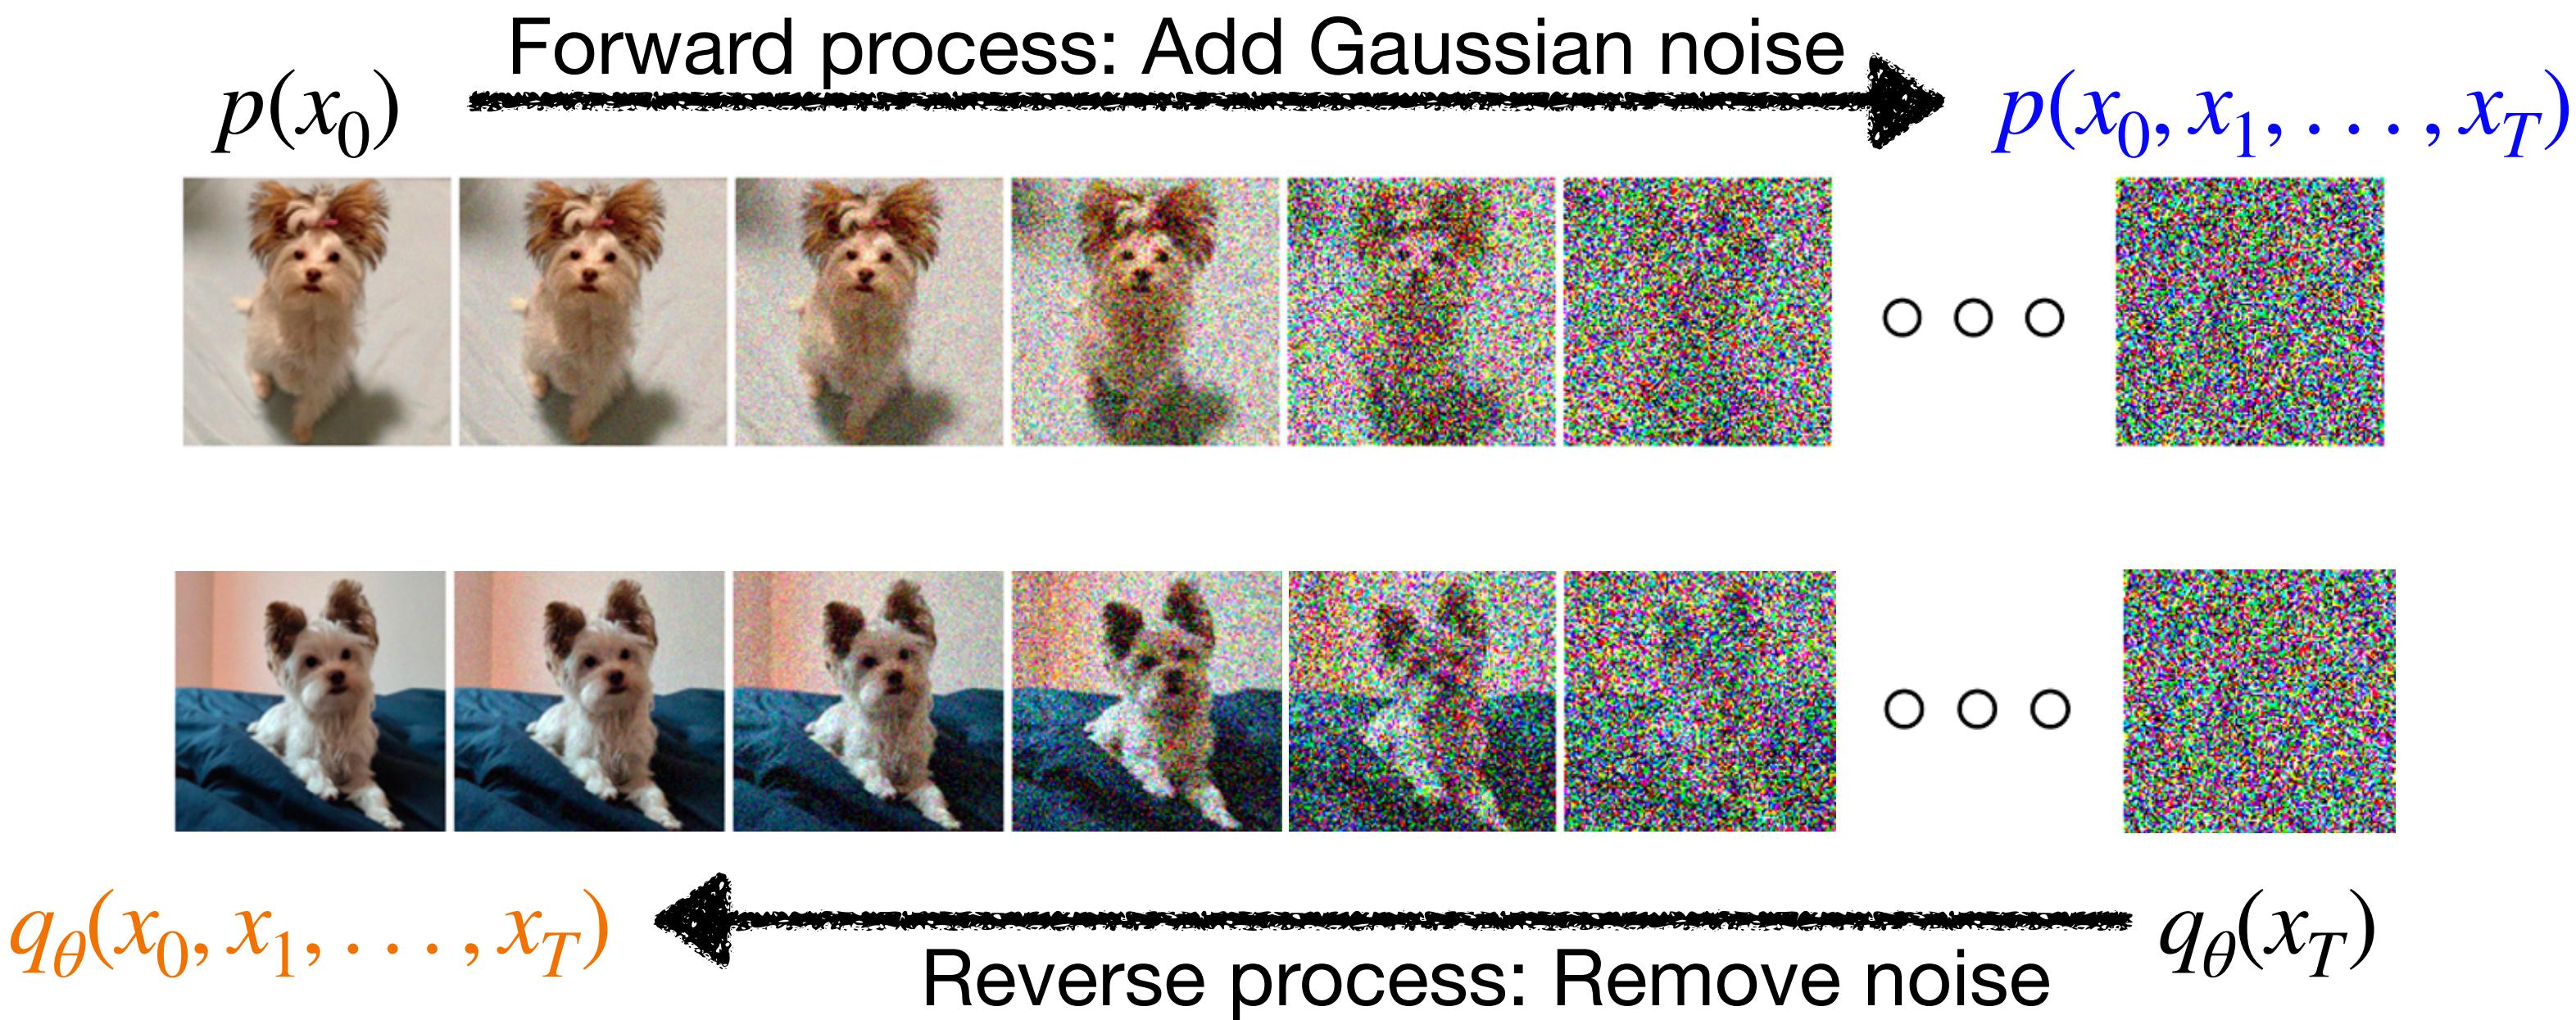
Why is the density evaluation not exact?

- In overview: Diffusion Models
- Density evaluation only exact for $T \rightarrow \infty$
 $\rightarrow q_\theta(x_0, x_1, \dots, x_T)$ for e.g. $T = 10$ is only approximation!

X No/approximate evaluation of the probability

Diffusion Models: 3 Take-Aways

1. Diffusion Model = Model that learns how to remove noise



2. Parametrize mean and variance of Gaussian by DNN

3. Density Evaluation is approximation

How to normalize? - Three Model Families

1. Approximate Z_θ

- Energy-based models

X Inaccurate probability

2. Use restricted NN architectures

- Variational autoencoders
- Autoregressive models
- **Normalizing Flows** ✓

X Limited flexibility of the model

3. Model the generative process only

- GANs
- **Diffusion models** ✓
- Score matching

X No/approximate evaluation of the probability

Recommended References for the Interested

- Energy-based models:
 - Song and Kingma, “[How to train your Energy-Based Model](#)”, 2021
- Normalizing Flows:
 - Rezende and Mohamed, “[Variational Inference with Normalizing Flows](#)”, 2016
 - Kobyzev, et al., “[Normalizing Flows: An Introduction and Review of Current Methods](#)”, 2021
- Diffusion Models:
 - Yang, et al. “[Diffusion Models: A Comprehensive Survey of Methods and Applications](#)”, 2024, under review
- Score-Matching:
 - Song, et al., “[Score-Based Generative Modeling through Stochastic Differential Equations](#)”, 2021
- Talk by Yang Song on “[Diffusion and Score-Based Generative Models](#)”, 2022

Thank you!
Do you have any questions?